

# NOVA SCHOOL OF SCIENCE AND TECHNOLOGY



## CONCURRENCY AND PARALLELISM

MASTER IN COMPUTER SCIENCE AND ENGINEERING

2020/2021 - 2<sup>ND</sup> SEMESTER

---

# Lecture and Lab Notes

---

### Author(s):

- Rúben BARREIRO - *r.barreiro@campus.fct.unl.pt*

Last updated: March 23, 2021

# Lecture 1A - Course Administrivia (March 15, 2021)

## 1A.1. Basic Information of Lecturers

The *professor* responsible for the *lectures* and *labs* is:

- Prof. *João* LOURENÇO - [joao.lourenco@fct.unl.pt](mailto:joao.lourenco@fct.unl.pt):
  - *Office Location*:
    - \* *Department of Informatics*;
    - \* Building II - Room P2/9, Ext. 10740;

## 1A.2. Discussion Forums

There will be available some *discussion forums* such as the following ones:

- *Piazza*:
  - [piazza.com/fct.unl.pt/spring2021/cp11158/home](https://piazza.com/fct.unl.pt/spring2021/cp11158/home);

## 1A.3. Main Bibliography

The *main bibliography* is the following:

- **Structured Parallel Programming: Patterns for Efficient Computation:**
  - *Michael* MCCOOL, *Arch* ROBINSON and *James* REINDERS;
  - *Morgan Kaufmann*, 2012;
  - ISBN: 978-0-12-415993-8;
  - *Click here to download*;
- **Patterns for Parallel Programming:**
  - *Tim* MATTSON, *Beverly* SANDERS and *Berna* MASSINGILL;
  - *Addison-Wesley*, 2014;
  - ISBN: 0-321-22811-1;
  - *Click here to download*;

- **Concurrent Programming: Algorithms, Principles, and Foundations:**

- *Michael RAYNAL*;
- *Springer-Verlag Berlin Heidelberg*, 2013;
- *ISBN: 978-3-642-32026-2*;
- *Click here to download*;

## 1A.4. Additional Bibliography

The *additional bibliography* is the following:

- **Programming Concurrency on the JVM: Mastering Synchronization, STM, and Actors:**

- *Venkat SUHRAMANIAM*;
- *Pragmatic Bookshelf*, 2011;
- *ISBN: 978-1-934356-76-0*;
- *Click here to download*;

- **The Art of Multiprocessor Programming:**

- *Maurice HERLIHY, Nir SHAVIT, Victor LUCHANGCO and Michael SPEAR*;
- *Morgan Kaufman*, 2021;
- *ISBN: 978-0-12-415950-1*;
- *Click here to download*;

- **Shared-Memory Synchronization:**

- *Michael SCOTT*;
- *Morgan & Claypool*, 2013;
- *ISBN: 978-1-608-45956-8*;
- *Click here to download*;

- **Principles of Concurrent and Distributed Programming:**

- *Michael BEN-ARI;*
- *Pearson, 2006;*
- *ISBN: 978-0-321-31283-9;*
- *Click here to download;*

## 1A.5. Other Bibliography

Other *recommended bibliography* is the following:

- **Pro Git:**

- *Scott CHACON and Scott CHACON;*
- *Apress, 2014;*
- *ISBN: 978-1-4842-0076-6;*
- *Click here to download;*

## 1A.6. Syllabus

The structure of the *course* is described by the following enumerated topics:

1. **Parallel Architectures:**

- *Flynn's Taxonomy;*
- *Performance Theory (including Amdahl's and Gustafson's Laws);*

2. **Parallel Programming:**

- *The spectrum of high-demanding computational problems;*
- *Regular and irregular problems;*
- *Strategies for problem decomposition and their mapping to programming patterns;*
- *The transactional and Map-Reduce models;*

### 3. Concurrency Control and Synchronization:

- *Competition and Collaboration;*
- *Atomicity;*
- *Linearization;*
- *Monitors;*
- *Locks*
- *Semaphores;*
- *Barriers;*
- *Producer-Consumer;*
- *Multi-Reader Single-Writer Locks;*
- *Futures;*
- *Concurrency in Practice in Java and C;*

### 4. Safety and Liveness:

- *Safety vs. Liveness;*
- *Progress;*
- *Deadlock;*
- *Deadlock Prevention, Avoidance, Detection, and Recovery;*
- *Livelock;*
- *Livelock Avoidance;*
- *Priority Inversion;*
- *Priority Inheritance;*
- *Lock-Free Algorithms;*

### 5. The Transactional Model:

- *Composite Operations;*
- *Transactions (Serializability);*
- *Optimistic Concurrency Control (OCC);*
- *Transactional Memory;*

## 6. Concurrency without Shared Data:

- *Active Objects*;
- *Message Passing*;
- *Actors*;

## 1A.7. Evaluation

- 60% - **2 Tests** (*Individual* and *Online*) [average  $\geq 8.5$  points];
- 40% - **1 Project** (Groups of 3 Students) [grade  $\geq 8.5$  points];
- 3% - Participation in *Classes' Life Cycle*:
  - *Lectures*;
  - *Labs*;
  - *Piazza*;

# Lecture 1B - Parallel Programming Models and Architectures (March 15, 2021)

## 1B.1. Introduction

All *Classical Computers* are now **parallel**, specially, the modern ones, which support **parallelism** in *hardware*, through, at least, one *parallel feature*:

- *Vector Instructions*;
- *Multi-threaded Cores*;
- ***Multicore Processors***;
- *Multiple Processors*;
- *Graphics Engines*;
- *Parallel Co-Processors*;

This statement does not apply only to *supercomputers*, but even the *smallest* computers, such as *phones*, support many of these features, where, it is necessary to use explicit ***parallel programming*** to get the most of them.

The *automatic approaches* to *parallelizing serial code* ***cannot*** deal with the shifts in the *algorithm's structures* required for effective *parallelization*.

A *programmer*, in a *modern computing environment*, should not just take advantage of *processors* with *multi-cores*, but, must be able to write ***scalable applications***, that can take advantage of **any** amount of ***parallel hardware***.

The feature of ***scaling*** requires attention to many factors:

- *Minimization of Data Movement*;
- *Serial Bottlenecks* (including *Locking*);
- Other forms of *Overhead*;

Some ***Parallel Patterns*** can help with this factors, but ultimately, is a responsibility of the *software developer* to produce a ***good algorithm design***.

In the recent years, it were also made some progress in other *paradigm* of *computing*, known as, ***Quantum Computing***, which is also, *inherently parallel*, taking advantage of the **quantum parallelism** offered on the *atomic* and *sub-atomic* levels of ***elementary particles***, such as, **photons**.

## 1B.2. Think Parallel

The **Parallelism** is an *intuitive* and *common human experience*, where, *programmers* naturally accept the concept of *parallel work* via a *group of workers*, often with *specializations*, in order to achieve a *major goal*.

Some *important* and *relevant* notions are:

- **Serialization:**
  - Act of putting some set of operations into a specific order;
- **Serial Semantics:**
  - *Semantics* used even though the *hardware* was naturally *parallel*;
- **Serial Illusion:**
  - *Mental model* of a *computer* executing operations, *sequentially*;
  - Has the problem of programmers came to depend on it too much;

The **Serialization** has its own *benefits*, such as, if someone reads a piece of *serial code* from top to bottom, will also be able to understand the *temporal order* of operations from the structure of the *source code*.

It helps that modern *programming languages* have evolved to use some *structured control flow* to emphasize this aspect of *serial semantics*.

Unless it was intentionally injected *randomness*, the *serial programs* are **deterministic**, doing always the same operations in the same order, giving the *same answer*, every time you run them with the *same inputs*.

The **Determinism** is useful for *debugging*, *verification*, and *testing*, but the *deterministic behavior* is not a natural characteristic of *parallel programs*.

Generally speaking, the *timing* of *task execution* in *parallel programs*, in particular the *relative timing*, is often **Non-Deterministic**.

Given that *parallelism* is necessary for *performance*, it would be useful to find an effective approach to *parallel programming* that retains as many of the benefits of *serialization* as possible, yet is also similar to existing practice.

Some of the known **Parallel Patterns** provide *structure* but they can also avoid the existing **Non-Determinism**, with a few easily visible exceptions where it is *unavoidable* or *necessary* for a better *performance*.

When eliminating *unnecessary serialization*, leading to *poor performance*, the current *programming tools* may have many **serial traps** built into them.



The ***serial traps*** are constructs that make, often *unnecessary*, their *serial assumptions* and they can also exist in the *design* of *algorithms* and in the abstractions used, to estimate the respective *complexity* and *performance*.

Often, there are two main steps to *think* and *program* in *parallel*:

1. Learning to recognize ***serial traps***;
2. Programming in terms of ***Parallel Patterns*** that capture their best practices and using *efficient* implementations of these *patterns*;

The most difficult part of learning to *program* in *parallel* is *recognizing* and *avoiding* ***serial traps*** (i.e., *assumptions* of *serial ordering*).

These assumptions are so *common* that, their existence is unnoticed and most of the *programming languages* *unnecessarily* overconstrain the *execution order*, making the *parallel execution* of their *programs* to be difficult.

*Parallelizing serial programs* written using *iteration constructs*, requires to be *recognized* the *semantics* used and *convert* them to the appropriate *parallel structure* and even better would be to *design programs* using them.

### 1B.3. Vocabulary and Notation

The two fundamental components of algorithms are ***tasks*** and ***data***, where a ***task*** operates on ***data***, either *modifying* it in place or *creating* new ***data***.

In a *parallel computation*, the *multiple tasks* need to be *managed* and *coordinated*, where, in particular, ***dependencies*** between ***tasks*** need to be *respected*, regarding their particular *ordering* of *execution*.

The ***dependencies*** are often but not always associated with the transfer of ***data***, between ***tasks*** and can be categorised as the following two kinds:

- ***Data Dependency***:
  - ***Tasks*** that *cannot* be executed, before some of its required ***data*** be ready, which can be generated by other ***tasks***;
- ***Control Dependency***:
  - *Events* or *side effects*, such as, the ones happening due to some *I/O operations*, which need to be *ordered*, in time;

For the *task management*, there is the ***fork-join*** pattern, where can exist:

- **Fork Points:**

- New *serial control flows* are created, by *splitting* an existing *serial control flow*, previously executed;

- **Join Points:**

- Two *separate* executed *serial control flows* which are *synchronized*, by *merging* them together in one *serial control flow*;

In a *single serial control flow*, the ***tasks*** are *ordered*, according to the *serial semantics*, and due to the implicit *serial control flow* before and after these *points*, the ***control dependencies*** are also needed between ***fork*** and ***join points*** and the ***tasks*** that *precede* and *follow* them, respectively.

## 1B.4. Simple Approaches to Parallelization

Some simple approaches for *parallelization* are:

- ***Distributed Memory***:

- *Parallel system* based on *processors* ***linked*** with a *fast network*;
- The *processors* ***communicate*** with each other, via *messages*;

- ***Owner Computes***:

- *Distribute* the ***data*** *elements* to *processors*;
- Each *processor* *updates* its own ***data*** *elements*;

- ***Single Program, Multiple Data (SPMD)***:

- All *machines* run the ***same*** *program* on *independent data*;
- The *dominant* form of *parallel computing*;

For ***Shared Memory***, the *latency* of *communication* between the *processors* is *low*, while, for the ***Distributed Systems***, occurs the opposite, with *high latencies* on the *communications* between the *processors*, where this *drawback* can be solved, by keeping some ***data*** *elements* of other *processors* *locally*, with what is called ***ghost cells***, *decreasing* a *latency* of its *communications*.

## 1B.5. Flynn's Characterization

The *parallel processors* can be divided into categories based on whether they have *multiple flows of control*, *multiple streams of data*, or on both:

- **Single Instruction, Single Data (SISD):**
  - *Standard non-parallel processor*, often referred as *scalar processor*;
  - The *individual performance* of *scalar processing* is important, since if it is slow, it can end up *dominating* the *global performance*;
- **Single Instruction, Multiple Data (SIMD):**
  - *Single task* executing simultaneously on *multiple elements* of *data*;
  - The number of *data elements* in a *SIMD's operation* can vary from a small number, such as the 4 to 16 elements in short *vector instructions*, to thousands, as in *streaming vector processors*;
  - The *SIMD processors* are also known as *array processors*, since consist on an *array of functional units* with a *shared controller*;
- **Multiple Instruction, Single Data (MISD):**
  - Not particularly useful and is not used;
- **Multiple Instruction, Multiple Data (MIMD):**
  - Separated *streams* of ***instructions*** and ***tasks***, each one with its own *flow of control*, operating on separated ***data*** *elements*;
  - Defines *multiple cores* in a *single processor*, *multiple processors* in a *single computer*, and *multiple computers* in a *single cluster*;
  - Represents a ***heterogeneous computer***, as *multiple processors* using *different architectures*, present in the *same computer system*;
  - An example would be a *host processor* and a *co-processor* with *different* sets of ***instructions*** and ***tasks***;

## 1B.7. Software Taxonomies

Some commonly used *Software Taxonomies* are the following ones:

- **Data Parallel (SIMD):**
  - *Parallelism* that is a result of *identical operations* being applied *concurrently* on *different data elements* (e.g., *matrix algorithms*);
  - Difficult to apply it to some *complex problems*;
- **Single Program, Multiple Data (SPMD):**
  - A *single application* is executing across *multiple processes/threads* on a *MIMD architecture*;
  - Most processes execute the same code but do not work in *lock-step*;
  - *Dominant* form of *parallel programming*;

## 1B.8. Shared Memory (SM)

- **Attributes:**
  - Has *global memory space*;
  - A *processor* use its own *cache* for a part of *global memory space*;
  - The *consistency* of the *cache* is maintained by *hardware*;
- **Advantages:**
  - *User-friendly programming* techniques (*OpenMP* and *OpenACC*);
  - Has *low latency* for *data sharing* between **tasks**;
- **Disadvantages:**
  - Has *global memory space-to-CPU* path may be a *bottleneck*;
  - Has *Non-Uniform Memory Access (NUMA)*;
  - The *programmer* is responsible for the details of *synchronization*;

## 1B.9. Distributed Memory (DM)

- **Attributes:**
  - The *memory* is *shared* amongst *processors* via *message passing*;
- **Advantages:**
  - The *memory scales* based on the number of *processors*;
  - The access to a *processor's* own *memory* is *fast*;
  - Is *cost effective*;
- **Disadvantages:**
  - Is *error prone*;
  - The *programmer* is responsible for the details of *communication*;
  - The *complex data structures* may be difficult to *distribute*;

## 1B.10. Software/Hardware Models

The *Software* and *Hardware* models ***do not*** need to match:

- **Distributed Memory** software on **Shared Memory** hardware:
  - **Message Passing Interface (MPI)**, which is designed for **Distributed Memory (DM)** hardware but is also available on **Shared Memory (SM)** systems;
- **Shared Memory** software on **Distributed Memory** hardware:
  - **Remote Memory Access (RMA)** included within **MPI-3**;
  - **Partitioned Global Address Space (PGAS)** languages:
    - \* **Unified Parallel C** (extension to **ISO C 99**);
    - \* **Coarray Fortran** (Fortran 2008);

## 1B.11. Difficulties of Parallelization

There are some common difficulties on *Parallelization*, such as:

- The *serialization* causes *bottlenecks*;
- The *workload* is not *distributed*;
- The *debugging* is *hard*;
- The *serial* approach **does not parallelize**;

## 1B.12. Performance

Often, one of the most common *serial traps* is the habit of discussing the *performance* of the *algorithms*, focusing only on the *minimization* of the *total amount* of *computational work* to be done, in *parallel programming*.

From this common habit, comes two main problems, such as:

- The *computation* itself, may not be the *bottleneck*, since, the *access* to *memory* or its *communication* may constrain the *performance*;
- The *potential* for *scaling performance* on a *parallel computer* is always constrained by the *algorithm's span*, which is the *execution time* needed for the longest *chain of tasks* that, must be performed *sequentially*;

# Lab 1 - $\pi$ Approximation by Monte Carlo Method

(March 18, 2021)