


INTRODUCTION

Parallelism and Concurrency: System and Environment

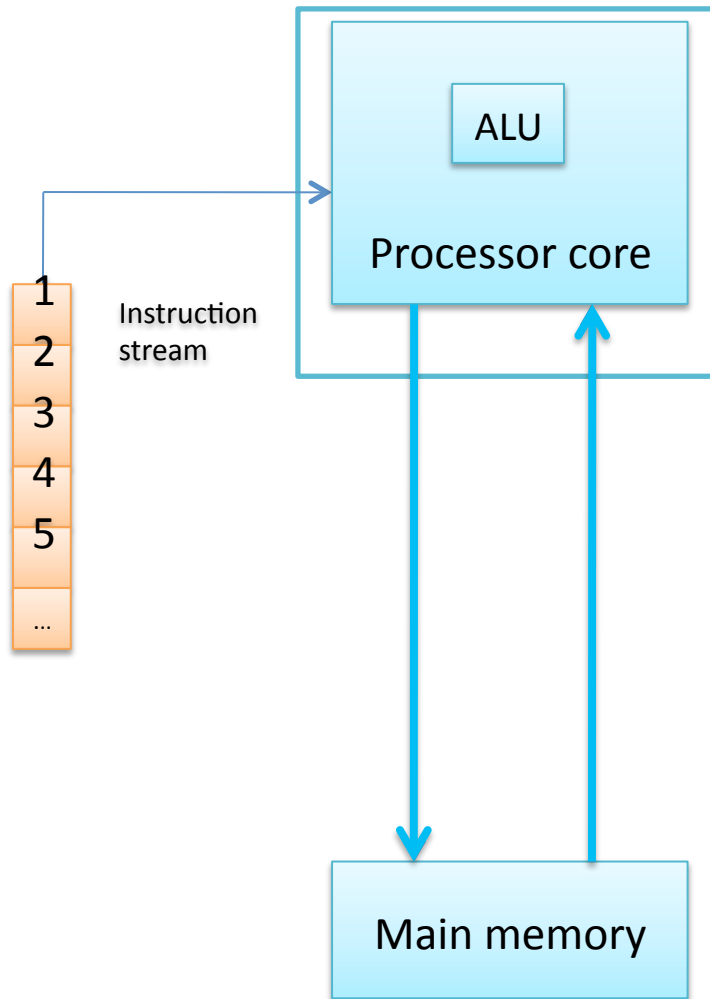
- *Parallelism*: exploit system resources to speed up computation
- *Concurrency*: respond quickly/ properly to events
 - from the environment
 - from other parts of system



Concurrency and Parallelism

- Why should you care?

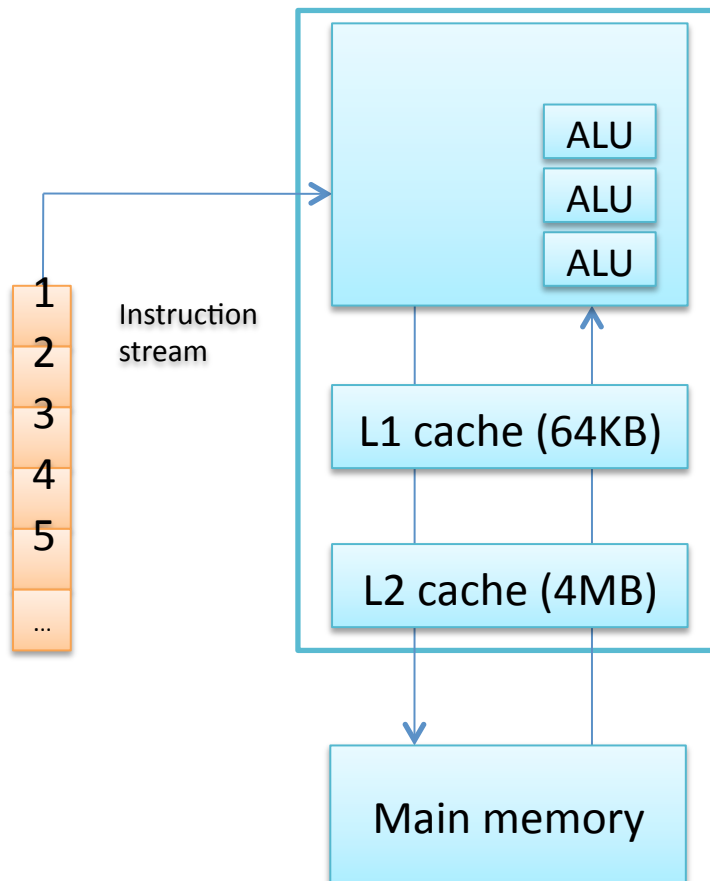
A simple microprocessor model ~ 1985




- Single h/w thread
- Instructions execute one after the other
- Memory access time ~ clock cycle time

ALU: arithmetic logic unit

FastFwd Two Decades (circa 2005): Power Hungry Superscalar with Caches



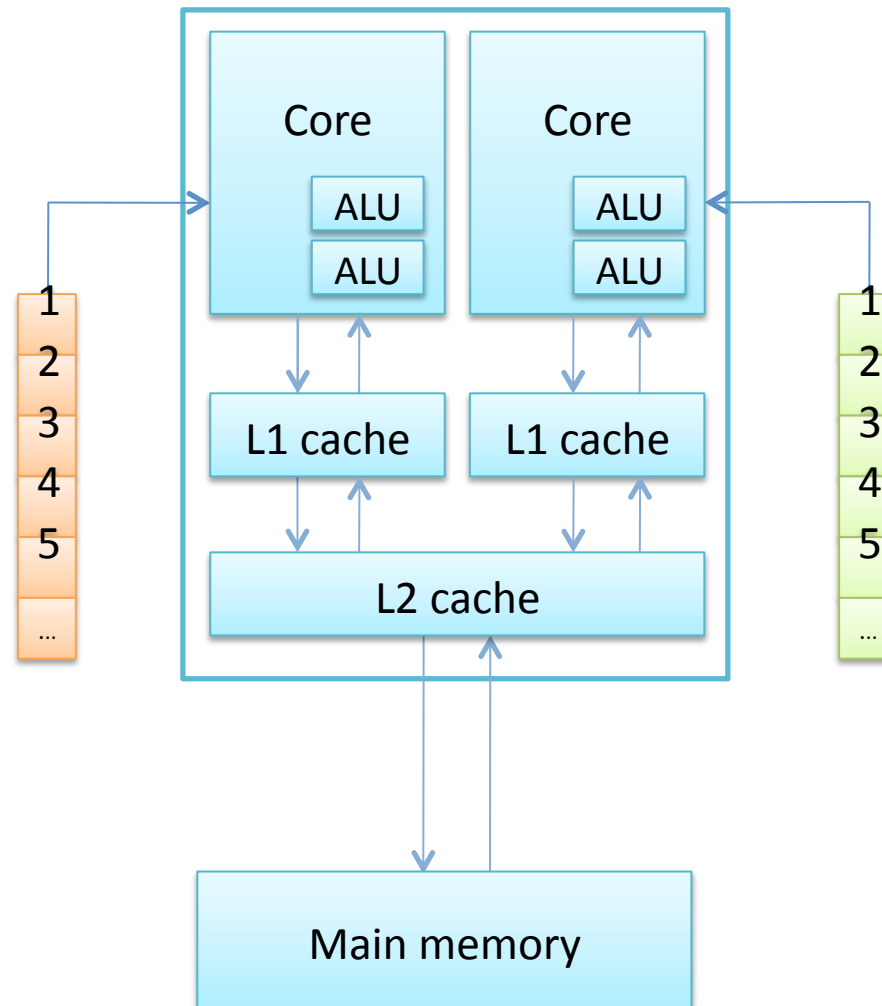
- Caching
- Pipelining
- Out-of-order execution
- Speculation
- ...



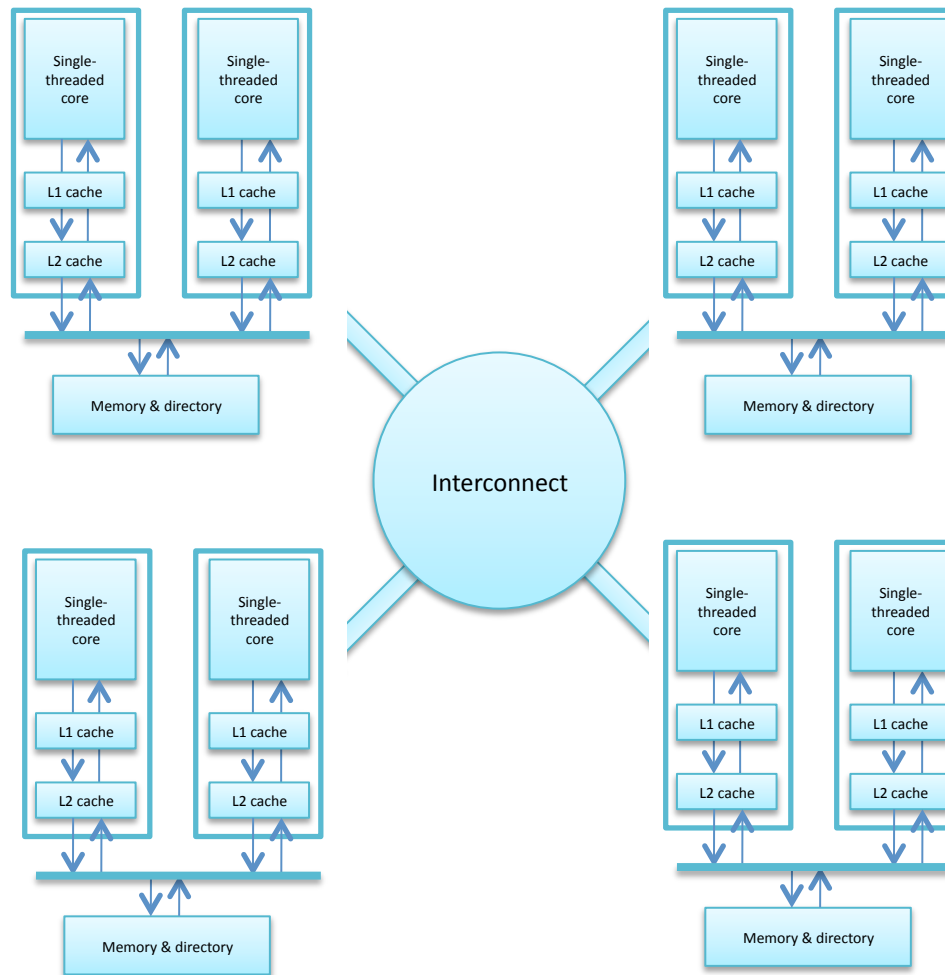
Power wall + Memory wall + ILP wall = **BRICK WALL**

- *Power wall*
 - we can't clock processors faster
- *Memory wall*
 - many workload's performance is dominated by memory access times
- *Instruction-level Parallelism (ILP) wall*
 - we can't find extra work to keep functional units busy while waiting for memory accesses

Multi-core h/w – common L2



NUMA multiprocessor



Technology Trends

- Increasing parallelism within a computer
 - Multi-core CPU
 - Graphical processing unit (GPU)
- Importance of power-efficient computing
 - Mobile phones, tablets
- Increasing disk capacity
 - We are awash in interesting *data*
 - Data-intensive problems require *parallel processing*

Technology Trends (2)

- Increasing *networks and network bandwidth*
 - Wireless, wimax, 3G, ...
 - Collection/delivery of massive datasets, plus
 - Real-time responsiveness to asynchronous events
- Increasing *number and variety of computers*
 - Smaller and smaller, and cheaper to build
 - Generating streams of asynchronous events

Application Areas

- Entertainment/games
- Finance
- Science
- Modeling of real-world
- Health care
- Telecommunication
- Data processing
- ...

Writing Concurrent and Parallel Code is HARD

- In general: difficult to do multiple things at the same time
- Performance considerations
- Correctness considerations

Performance Considerations

- Speedup
- Responsiveness
- Locality
- Load balance
- Coordination and synchronization

Correctness Concerns

- All those we have for sequential code
 - Assertions, invariants, contracts,
 - buffer overflows, null reference,
 - ...
- Plus those related to parallelism/concurrency
 - Data races, deadlocks, livelocks, ...
 - Memory coherence/consistency