

PARALLEL PERFORMANCE

Performance Considerations

- I parallelized my code and ...
- The code slower than the sequential version, or
- I don't get enough speedup
- What should I do?

Performance Considerations

- Algorithmic
- Fine-grained parallelism
- True contention
- False contention
- Other Bottlenecks

Algorithmic Bottlenecks


- Parallel algorithms might sometimes be completely different from their sequential counterparts
- Might need different design and implementation

Algorithmic Bottlenecks

- Parallel algorithms might sometimes be completely different from their sequential counterparts
- Might need different design and implementation
- Example: MergeSort (once again ☹)

Recall from Previous Lecture

- Merge was the bottleneck



Most Efficient Sequential Merge is not Parallelizable

```
Merge(int* a, int* b, int* result ){  
  
    while( <end condition> ) {  
        if(*a <= *b){  
            *result++ = *a++;  
        }  
        else {  
            *result++ = *b++;  
        }  
    }  
  
}
```

Parallel Merge Algorithm

- Merge two sorted arrays A and B using divide and conquer

Parallel Merge Algorithm

- Merge two sorted arrays A and B
 1. Let A be the larger array (else swap A and B)
 2. Let n be the size of A
 3. Split A into $A[0 \dots n/2 - 1]$, $A[n/2]$, $A[n/2 + 1 \dots n]$
 4. Do binary search to find smallest m such that
$$B[m] \geq A[n/2]$$
 5. Split B into $B[0 \dots m - 1]$, $B[m, \dots]$
 6. return Merge($A[0 \dots n/2 - 1]$, $B[0 \dots m - 1]$), $A[n/2]$, Merge($A[n/2 + 1 \dots n]$, $B[m, \dots]$)

Assignment 1 Extra Credit

- Implement the Parallel Merge algorithm and measure performance improvement with your work stealing queue implementation

Fine Grained Parallelism

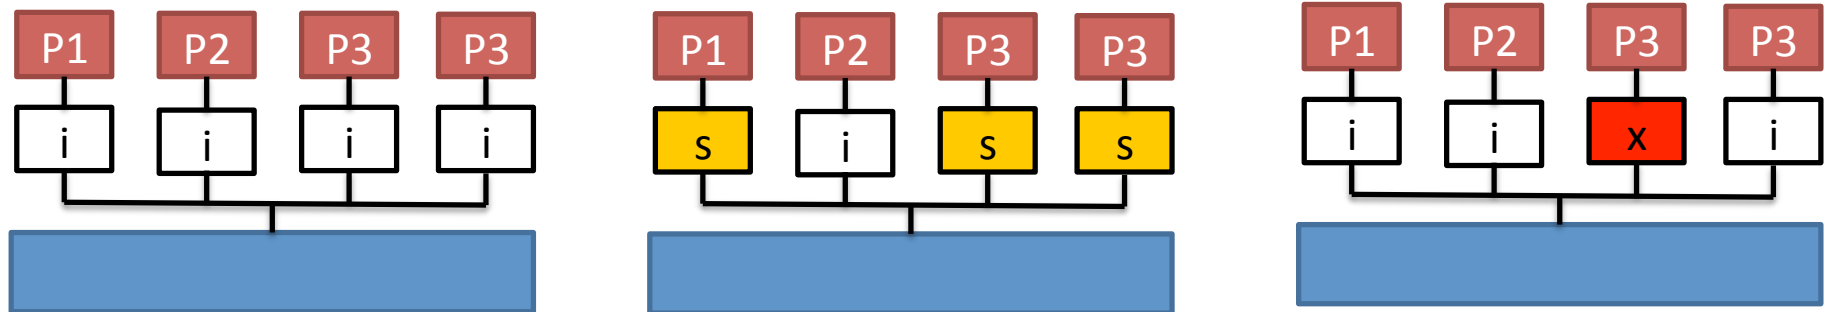
- Overheads of Tasks
 - Each Task uses some memory (not as much resources as threads, though)
 - Work stealing queue operations
- If the work done in each task is small, then the overheads might not justify the improvements in parallelism

False Sharing

Data Locality & Cache Behavior

- Performance of computation depends HUGELY on how well the cache is working
- Too many cache misses, if processors are “fighting” for the same cache lines
- Even if they don’t access the same data

Cache Coherence



- Each cacheline, on each processor, has one of these states:
 - i - invalid : not cached here
 - s - shared : cached, but immutable
 - x - exclusive: cached, and can be read or written
- State transitions require communication between caches (cache coherence protocol)
 - If a processor writes to a line, it removes it from all other caches

Ping-Pong & False Sharing

- Ping-Pong
 - If two processors both keep writing to the same location, cache line has to go back and forth
 - Very inefficient (lots of cache misses)
- False Sharing
 - Two processors writing to two different variables may happen to write to the same cacheline
 - If both variables are allocated on the same cache line
 - Get ping-pong effect as above, and horrible performance

False Sharing Example

```
void WithFalseSharing()
{
    Random rand1 = new Random(), rand2 = new Random();
    int[] results1 = new int[20000000],
        results2 = new int[20000000];
    Parallel.Invoke(
        () => {
            for (int i = 0; i < results1.Length; i++)
                results1[i] = rand1.Next();
        },
        () => {
            for (int i = 0; i < results2.Length; i++)
                results2[i] = rand2.Next();
        });
}
```

False Sharing Example

```
void WithFalseSharing()
{
    Random rand1 = new Random(), rand2 = new Random();
    int[] results1 = new int[20000000],
        results2 = new int[20000000];
    Parallel.Invoke(
        () => {
            for (int i = 0; i < results1.Length; i++)
                results1[i] = rand1.Next();
        },
        () => {
            for (int i = 0; i < results2.Length; i++)
                results2[i] = rand2.Next();
        });
}
```

rand1, rand2
are allocated
at same time
=>
likely on same
cache line.

Call to Next()
writes to the
random
object
=>
Ping-Pong
Effect

False Sharing, Eliminated?

```
void WithoutFalseSharing()
{
    int[] results1, results2;
    Parallel.Invoke(
        () => {
            Random rand1 = new Random();
            results1 = new int[20000000];
            for (int i = 0; i < results1.Length; i++)
                results1[i] = rand1.Next();
        },
        () => {
            Random rand2 = new Random();
            results2 = new int[20000000];
            for (int i = 0; i < results2.Length; i++)
                results2[i] = rand2.Next();
        });
}
```

rand1, rand2
are allocated
by different
tasks
=>
Not likely on
same cache
line.