

CORRECTNESS CRITERIA FOR CONCURRENCY & PARALLELISM

Contracts – Correctness for Sequential Code

Contracts – Correctness for Sequential Code

- Assertions
 - A predicate expected to hold at a particular program point
- Precondition
 - A predicate expected to hold at a function call
 - A failure can be blamed on the caller
- Postcondition
 - A predicate expected to hold at a function return
 - A failure can be blamed on the callee

Code Contracts for .Net

```
int Divide(int n, int d) {  
    return n/d;  
}
```

Preconditions using Requires

```
int Divide(int n, int d) {  
    Contract.Requires( 0 != d );  
    return n/d;  
}
```

Preconditions using Ensures

```
int Divide(int n, int d) {  
    Contract.Requires( 0 != d );  
    Contract.Ensures(  
        Contract.Result<int>() * d <= n &&  
        Contract.Result<int>() * d > n-d  
    );  
    return n/d;  
}
```

Example: Library APIs

vector::push_back

```
void push_back ( const T& x );
```

Add element at the end

Adds a new element at the end of the vector, after its current last element. The element is initialized to a copy of *x*.

This effectively increases the vector *size* by one, which causes a reallocation of the vector if its *size* was equal to the vector *capacity* before the call. Reallocations invalidate references and pointers.

Parameters

x

Value to be copied to the new element.

T is the first template parameter (the type of the elements stored in the vector).

Return value

none

If a reallocation happens, it is performed using `Allocator::allocate()`, which may throw `std::bad_alloc` if the allocation request does not succeed).

Example: System Call API

NAME

`connect` -- initiate a connection on a socket

SYNOPSIS

```
#include <sys/types.h>
#include <sys/socket.h>

int
connect(int socket, const struct sockaddr *address,
        socklen_t address_len);
```

DESCRIPTION

The parameter `socket` is a socket. If it is of type `SOCK_DGRAM`, this call specifies the peer with which the socket is to be associated; this address is that to which datagrams are to be sent, and the only address from which datagrams are to be received. If the socket is of type `SOCK_STREAM`, this call attempts to make a connection to another socket. The other socket is specified by `address`, which is an address in the communications space of the socket.

Each communications space interprets the `address` parameter in its own way. Generally, stream sockets may successfully `connect()` only once; datagram sockets may use `connect()` multiple times to change their association. Datagram sockets may dissolve the association by connecting to an invalid address, such as a null address or an address with the address family set to `AF_UNSPEC` (the error `EAFNOSUPPORT` will be harmlessly returned).


RETURN VALUES

Upon successful completion, a value of 0 is returned. Otherwise, a value of -1 is returned and the global integer variable `errno` is set to indicate the error.

Correctness for Concurrency & Parallelism

- Reuse contracts written for sequential code
- Relate correctness of concurrent/parallel executions to correctness of appropriate sequential executions

Coin Sorting Example



Use Contracts for Correctness Criteria

- SortCoins accepts a set of coins and returns a set of bad ones
- Parallelizing SortCoins should not change the contract

```
SortCoins(...) {  
  Contract.Requires(...);  
  Contract.Ensures(...);
```

Sequential
Implementation


```
}
```

```
SortCoins(...) {  
  Contract.Requires(...);  
  Contract.Ensures(...);
```

Parallel
Implementation

```
}
```


Parallelizing can sometimes produce correct but different outputs



Sequential
SortCoins

accept

reject



Parallel
SortCoins

accept

reject



The order of coins can change

Restrictive Contracts Can Limit Parallelism

- The order of coins returned by SortCoins might be different from the input order
- Do care about the total amount returned
- If the contracts enforce the ordering, resulting parallelization might be unacceptably slow
- Design interfaces in such a way that contracts are not restrictive

Strategies For Reasoning About Parallel Code

- How do we know if a parallel loop satisfies its contract?
 - Reasoning about parallel executions is hard

General Strategy:

1. Make sure that every parallel behavior is equal to some sequential behavior
2. Convince that the sequential behavior satisfies the contract

Independent Loops

- Let $m(1) \dots m(n)$ be the loop iterations
- Two iterations $m(j)$ and $m(k)$ ($j \neq k$) are *dependent* if they access the same memory location and at least one of them is a write.
- Loop iterations are *independent* if no two of them are dependent

Dependent Loops

- Dependencies need to be made explicit
 - No data races
- Reason that order of dependent operations don't matter
 - e.g. These operations are commutative and associative
 - Recall: reduce/scan

Determinism:

A New Correctness Criteria

- Pre and post conditions do two things
 - Specify how the function behaves sequentially
 - Enforce the same behavior when parallelized

```
int ComputeSum ( IEnumerable<int> input)
{
    Contract.Requires ( input != null);

    Contract.Ensures (
        Contract.Result<int>() == input.Sum(i => i)
    );

    //implementation
}
```

Determinism Contract

- Allows you to check parallel correctness
- Without having to specify the sequential contract
 - The output of the function does not depend on task interleavings for a given input

```
int ComputeSum ( IEnumerable<int> input)
{
    Contract.Requires ( input != null);

    IsDeterministic (
        Contract.Result<int>(), input );

    //implementation
}
```

Determinism Checking

```
Contract.IsDeterministic (  
    output, {input1, input2, ... })
```

Is same as saying

```
Contract.Ensures (  
    output == F ( input1, input2, ... ) )
```

For some deterministic function F

Very useful when specifying F is tedious

Determinism Checking

```
Contract.IsDeterministic (  
    output, {input1, input2, ... }, comp)
```

Is same as saying

```
Contract.Ensures (  
    comp(output, F ( input1, input2,  
... ) )
```

For some deterministic function F

Strategies for Checking Determinism

Concurrent Objects

- Can be called concurrently by many threads
- Examples
 - Work Stealing Queue

Concurrent Objects

- Can be called concurrently by many entities
- Examples
 - Work Stealing Queue
 - C Runtime library
 - Operating System
 - Data bases


Correctness Criteria

- Informally called “thread safety”
- What does “thread safety” mean to you?

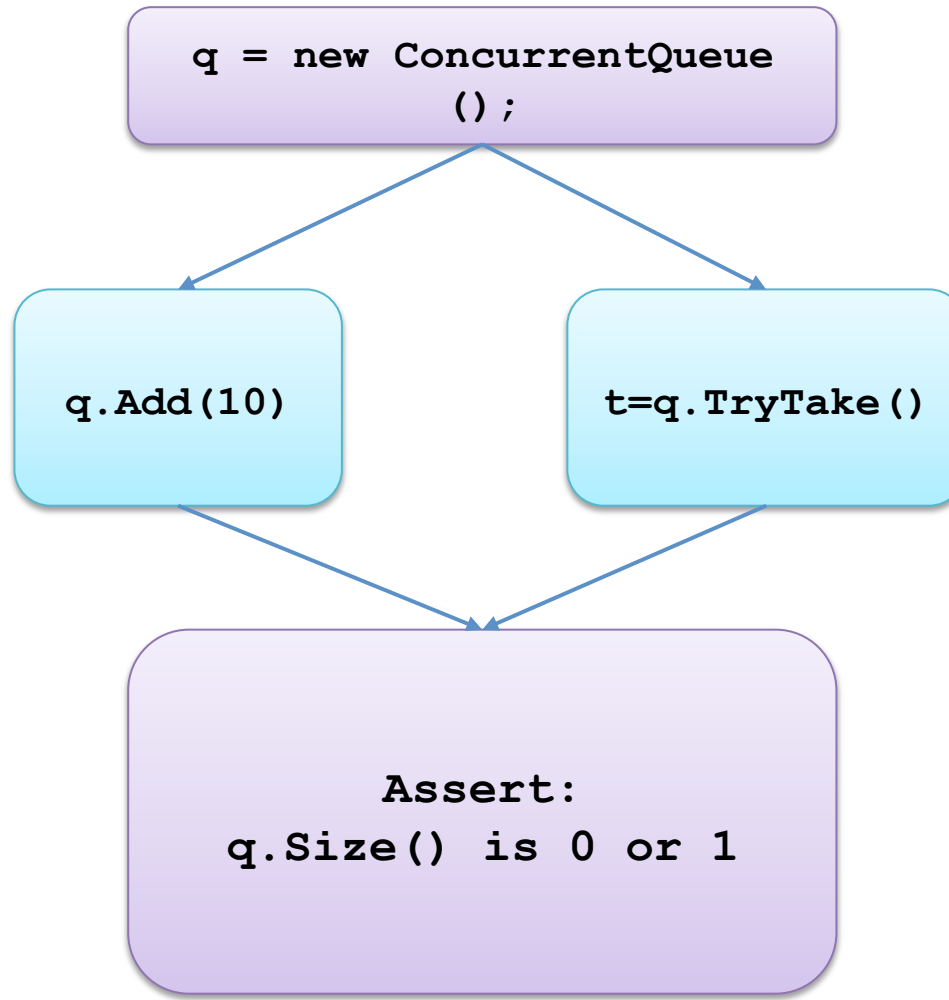
A Simple Concurrent Object

- Sequential Queue
 - Add(item)
 - TryTake() returns an item or “empty”
 - Size() returns # of items in queue
- Consider ConcurrentQueue and its relationship to Queue


Let's Write a Test



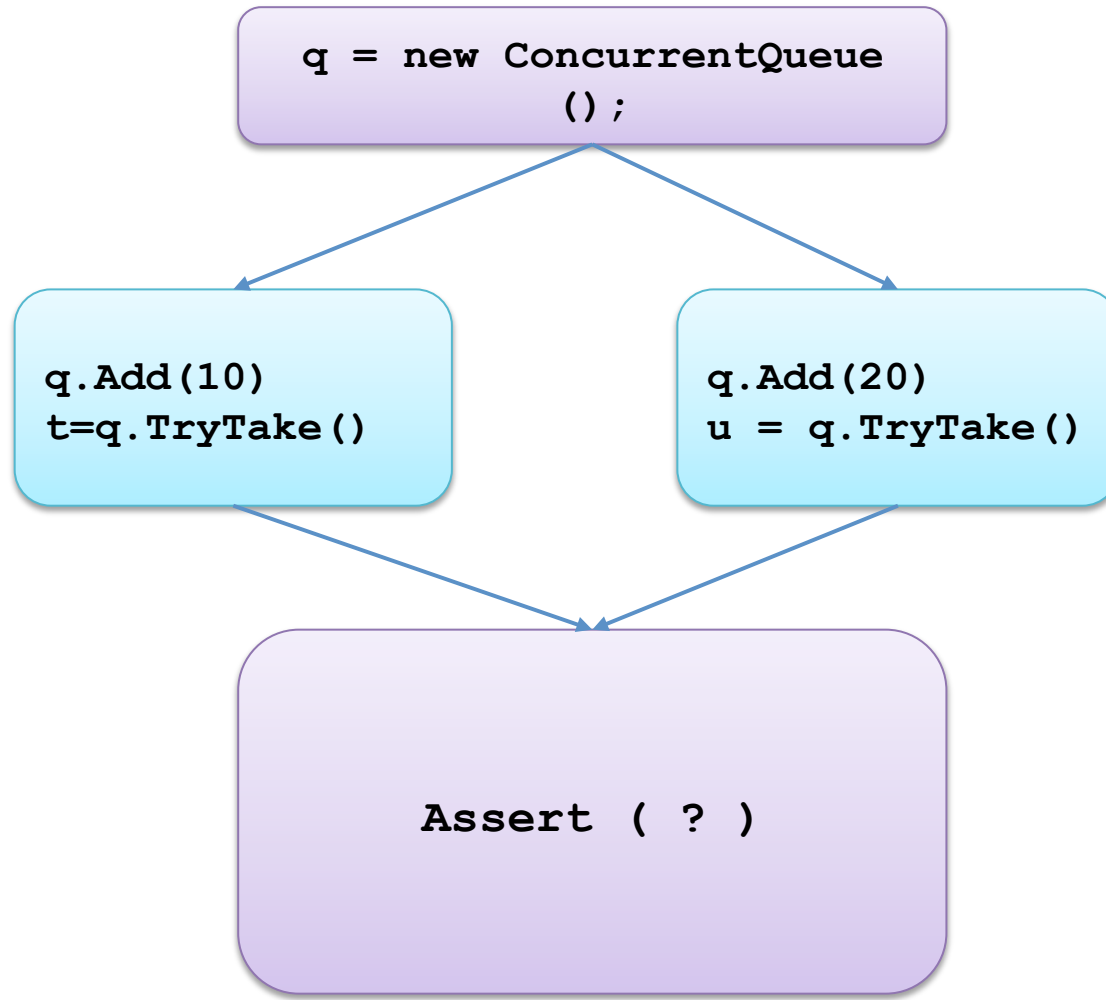
Let's Write a Test




Let's Write a Test



Let's Write a Test



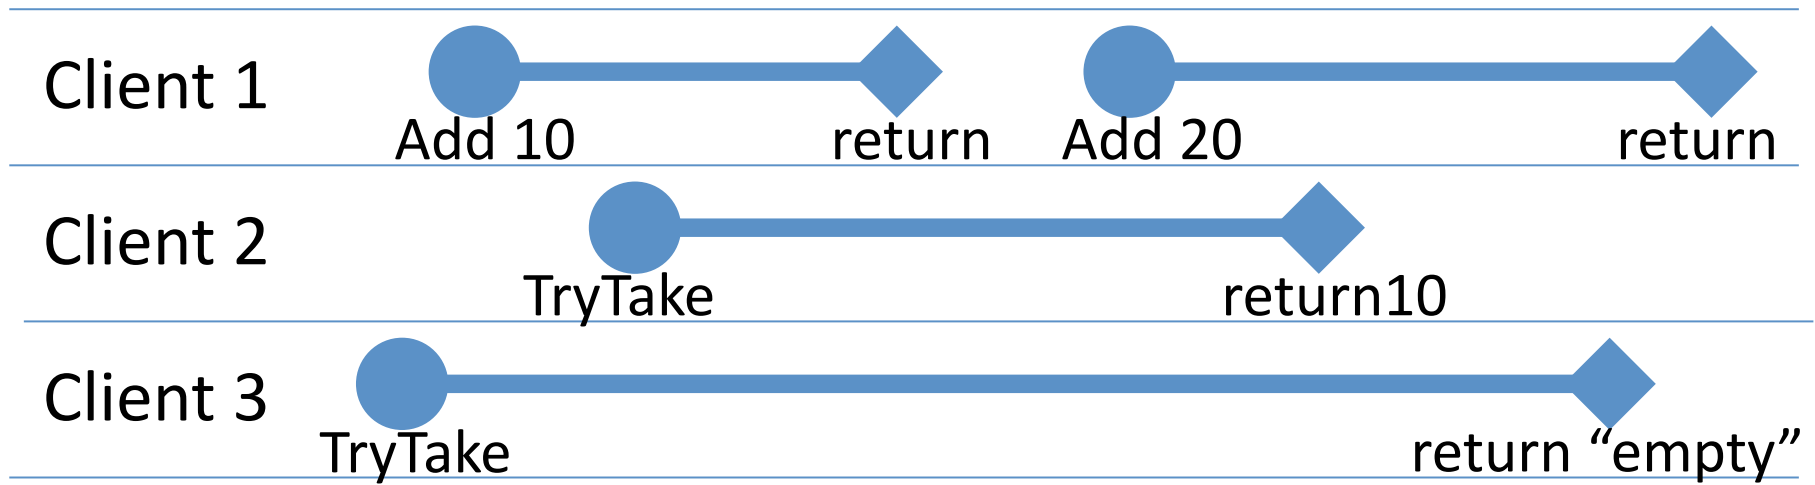
Let's Write a Test



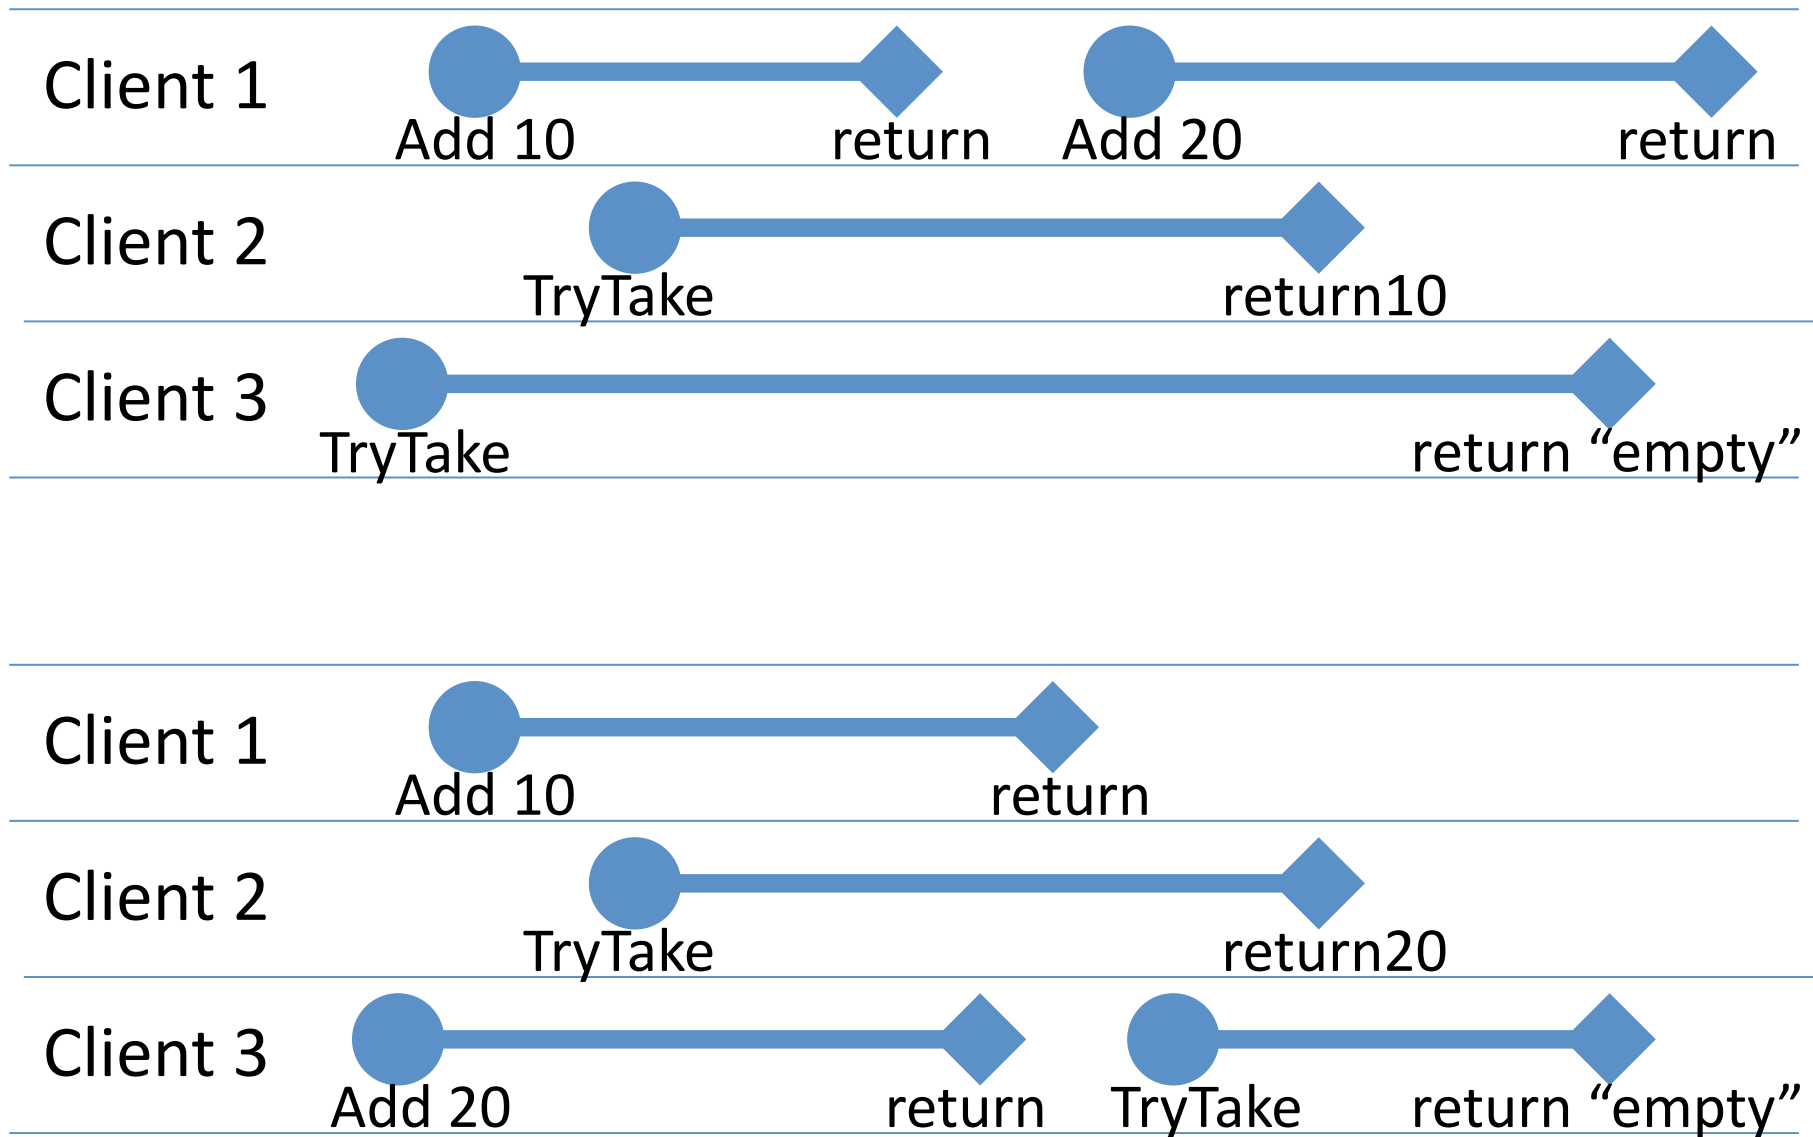
Linearizability

- The correctness notion closest to “thread safety”
- A concurrent component behaves ***as if*** only one thread can enter the component at a time

“Expected” Behavior?

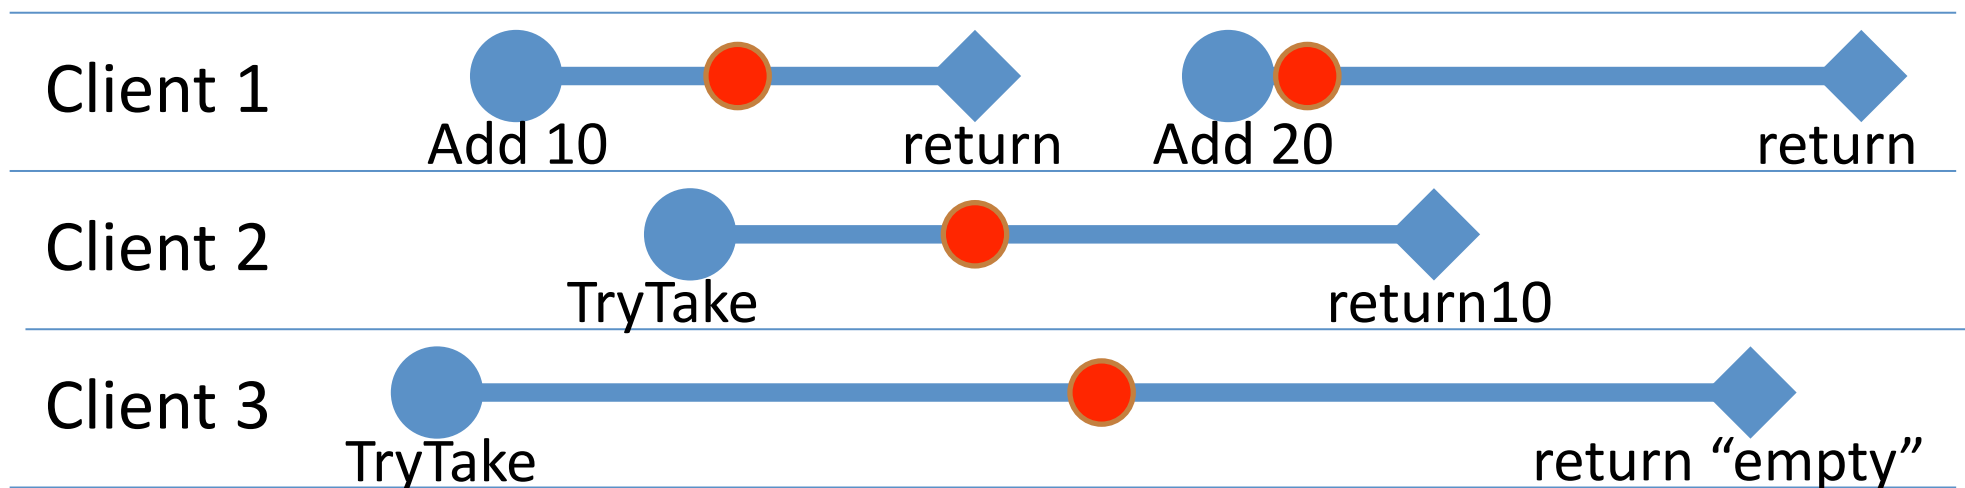


“Expected” Behavior?



Linearizability

- Component is *linearizable* if all operations
 - Appear to take effect atomically at a single temporal point
 - And that point is between the call and the return
- “As if the requests went to the queue one at a time”




Linearizability vs Serializability?

- Serializability
 - All operations (transactions) appear to take effect atomically at a single temporal point

Linearizability vs Serializability?


- Serializability
 - All operations (transactions) appear to take effect atomically at a single temporal point
- Linearizability
 - All operations to take effect atomically at a single temporal point
 - That point is between the call and return

Serializable behavior that is not Linearizable



- Linearizability assumes that there is a global observer that can observe that Thread 1 finished before Thread 2 started


Serializability does not compose



- The behavior of the blue queue and green queue are individually serializable
- But, together, the behavior is not serializable


Formalizing Linearizability

- Define the set of observables for each operation
 - Call operation: value of all the arguments
 - Return operation:
- An event:
 - Thread Id, Object Id, Call/Return, Operation, Observables



A Concurrent History


- Sequence of Events
 - $\langle T1, q, \text{Call}, \text{Add}, 10 \rangle$
 - $\langle T2, q, \text{Call}, \text{TryTake}, \text{void} \rangle$
 - $\langle T1, q, \text{Ret}, \text{Add}, \text{void} \rangle$
 - $\langle T2, q, \text{Ret}, \text{TryTake}, \text{"empty"} \rangle$



A Concurrent History

- Sequence of Events
 - $\langle T1, q, \text{Call}, \text{Add}, 10 \rangle$
 - $\langle T2, q, \text{Call}, \text{TryTake}, \text{void} \rangle$
 - $\langle T1, q, \text{Ret}, \text{Add}, \text{void} \rangle$
 - $\langle T2, q, \text{Ret}, \text{TryTake}, \text{"empty"} \rangle$


We will only focus on single object histories



A Concurrent History


- Sequence of Events
 - $\langle T1, q, \text{Call}, \text{Add}, 10 \rangle$
 - $\langle T2, q, \text{Call}, \text{TryTake}, \text{void} \rangle$
 - $\langle T1, q, \text{Ret}, \text{Add}, \text{void} \rangle$
 - $\langle T2, q, \text{Ret}, \text{TryTake}, \text{"empty"} \rangle$

Also, we will only focus on complete histories – every call has a return



A Serial History

- A concurrent history where every call is followed by its matching return
 - $\langle T1, q, \text{Call}, \text{Add}, 10 \rangle$
 - $\langle T1, q, \text{Ret}, \text{Add}, \text{void} \rangle$
 - $\langle T2, q, \text{Call}, \text{TryTake}, \text{void} \rangle$
 - $\langle T2, q, \text{Ret}, \text{TryTake}, \text{"empty"} \rangle$

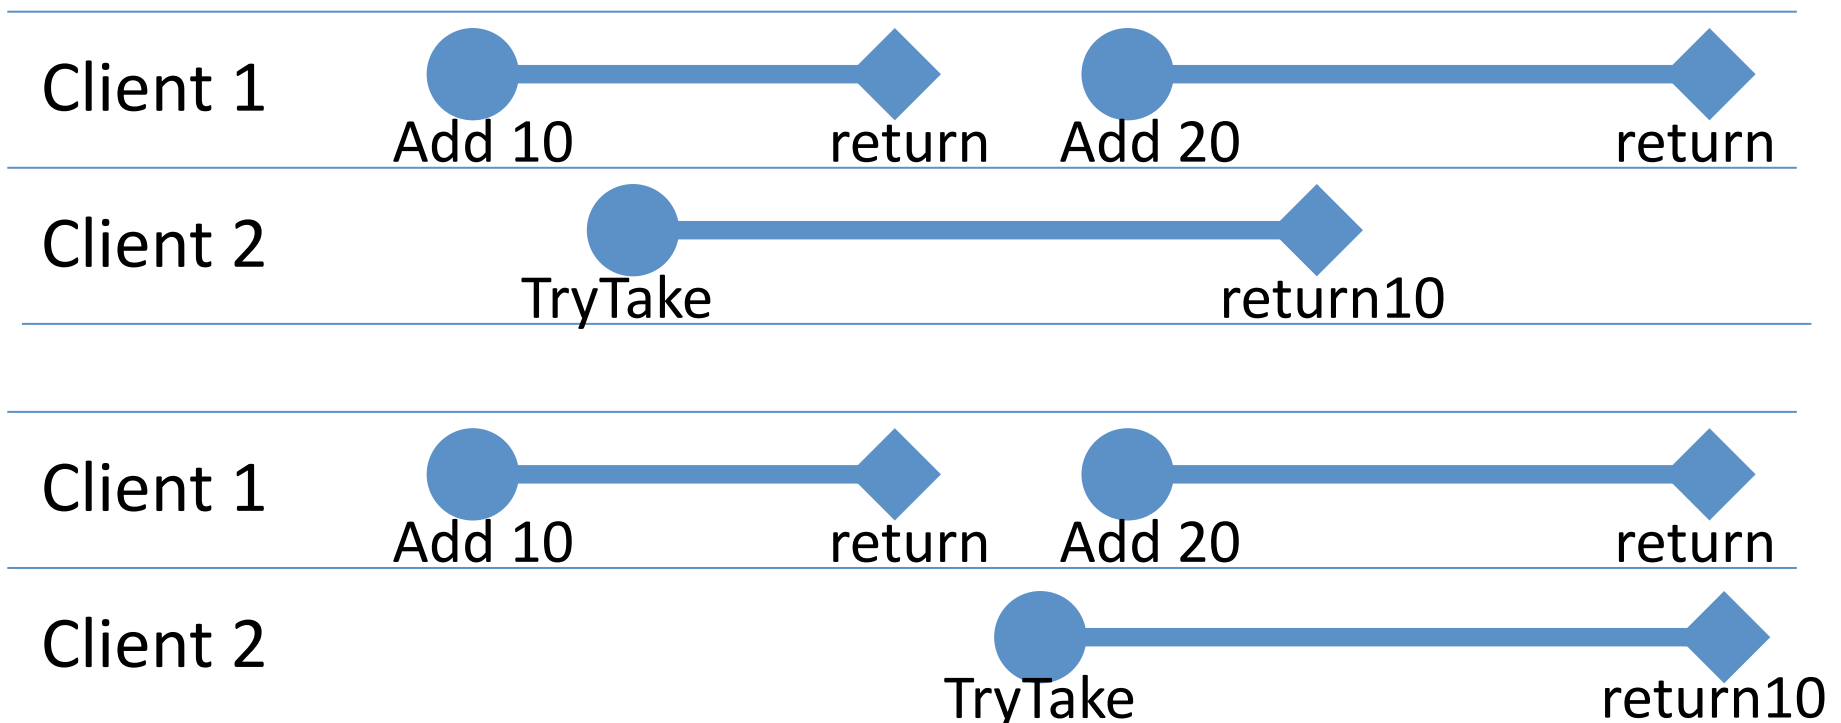


Sequential Specification of an Object

- The set of all serial histories define the sequential behavior of an object
- Assume we have a mechanism to enumerate this set and store the set in a database


Equivalent Histories

- Two concurrent histories are equivalent if
 - Each thread performs operations in the same order
 - And sees the same observations




Concurrent Operations in a History

- Two operations p and q are concurrent in a history if their duration overlap
 - $\neg (p.\text{ret} < q.\text{call} \vee q.\text{ret} < p.\text{call})$



Concurrent Operations in a History

- Two operations p and q are concurrent in a history if their duration overlap
 - $\neg (p.\text{ret} < q.\text{call} \vee q.\text{ret} < p.\text{call})$
- Non-Concurrent operations define a “performed-before” order



Linearizability

- A concurrent history is linearizable if it is equivalent to a (serial) history in the sequential specification,
- Such that all operations that are “performed before” in the concurrent history are also “performed before” in the serial history

