

# EFFICIENT DATA-RACE DETECTION

---

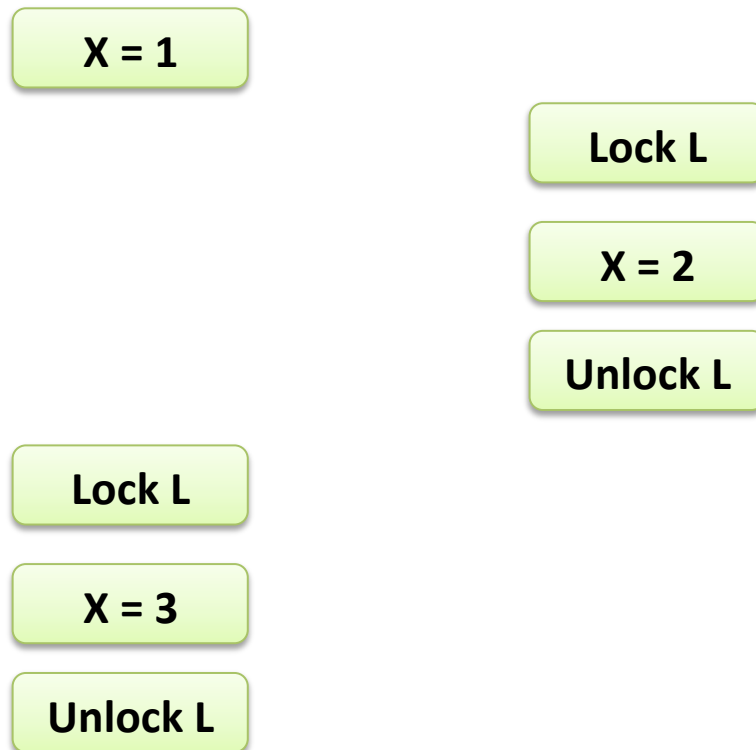
# Recap

- Happens-before based data-race detection
  - Shaz's guest lecture
  - Vector clock algorithm
- Lock-set based data-race detection
  - Eraser paper discussion

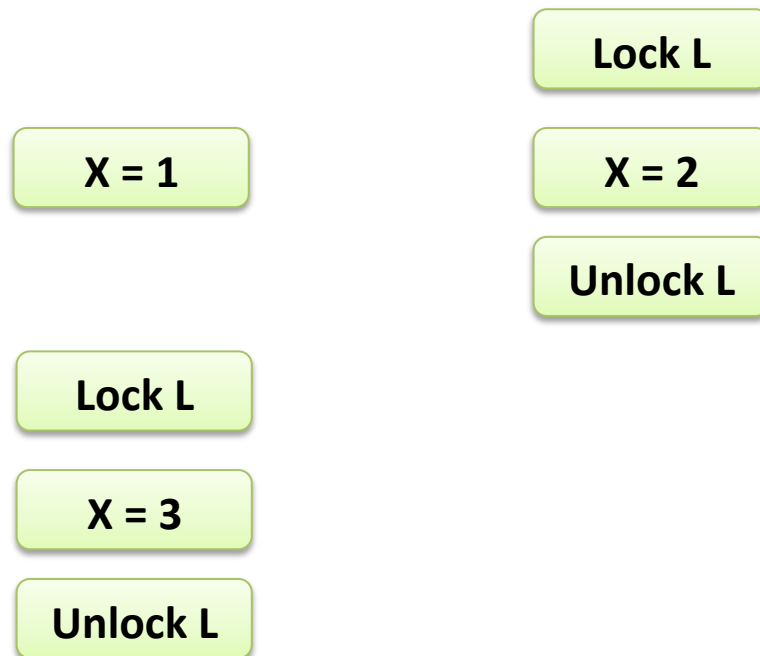
# New Definition of a Data Race

- Two instructions conflict if
  - They access the same memory location
  - At least one of them is a write
- A program contains a data race if
  - One can schedule the program on a multiprocessor
  - Such that conflicting instructions execute simultaneously

# New Definition vs Happens-Before



# New Definition vs Happens-Before



# New Definition vs Happens-Before

- True or False ?
- If a program has a data race as per the **new definition**
- Then it has a data race as per **the happens-before definition**

# New Definition vs Happens-Before

- True or False ?
- If a program has a data race as per the **new definition**
- Then it has a **happens-before** data race
- If a program has a **happens-before** data race
- Then it has a data race as per the **new definition**

# New Definition vs Lock Set

- True or False ?
- If a program has a data race as per the **new definition**
- Then it has a **lock-set based** data race
- If a program has a **lock-set based** data race
- Then it has a data race as per the **new definition**



# Challenges For Dynamic Data Race Detection

- Dynamically monitor memory accesses
  - Need instrumentation infrastructure
  - Performance overhead

# Challenges For Dynamic Data Race Detection

- Dynamically monitor memory accesses
  - Need instrumentation infrastructure
  - Performance overhead
- Understand synchronization mechanisms
  - Handle home-grown locks
  - Need to annotate the happens-before relationship

# Challenges For Dynamic Data Race Detection

- Dynamically monitor memory accesses
  - Need instrumentation infrastructure
  - Performance overhead
- Understand synchronization mechanisms
  - Handle home-grown locks
  - Need to annotate the happens-before relationship
- Maintain detection meta-data per variable, per synchronization object
  - Manage memory
  - Performance overhead

# DataCollider [OSDI '10]

- Lightweight
  - < 5% overhead
- Effective
  - Found data races in all applications we have run
- Easy to implement
  - The algorithm can be described in ~ 10 lines

# Design Choice #1

- Sample memory accesses
  - Detect data races on a small, random sample of memory accesses

# Design Choice #1

- Sample memory accesses
  - Detect data races on a small, random sample of memory accesses
- It is acceptable to miss races
  - Dynamic techniques cannot find all data races anyway
- A lightweight tool will have more coverage
  - Likely to be run on a lot more tests
- But data races are 'rare' events
  - Need to be careful when sampling

## Design Choice #2

- Perturb thread schedules
  - By inserting delays
  - Force the data race to happen, rather than infer from past memory/synchronization accesses

## Design Choice #2

- Perturb thread schedules

**X = 1**

**Lock L**

**X = 2**

**Unlock L**

**Lock L**

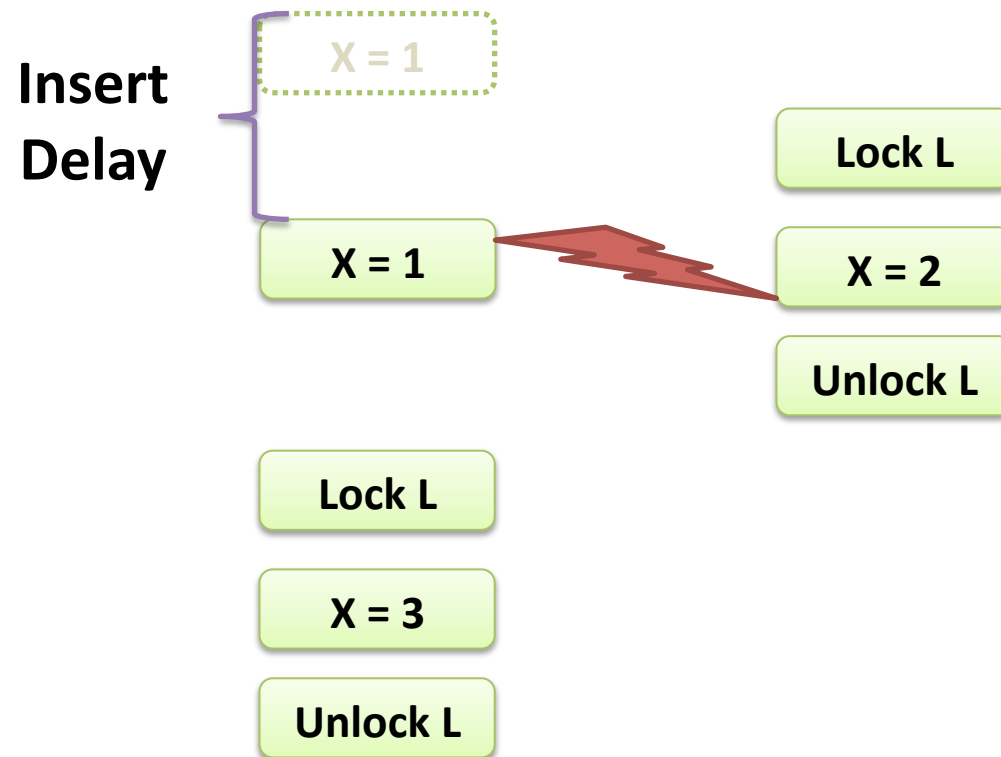
**X = 3**

**Unlock L**



## Design Choice #2

- Perturb thread schedules



# The Algorithm

```
At_Every_Memory_Access ( x ) {  
    if (not sampled) return;  
  
    delay for some time;  
  
    if another thread performed  
        a conflicting operation  
        during the delay  
        report data race;  
}
```

## Use Data Breakpoints for Detecting Conflicting Operations

```
At_Every_Memory_Access ( x ) {  
    if (not sampled) return;  
  
    rw = accessIsWrite ? RW : W;  
    SetDataBreakpoint ( x, rw );  
  
    Delay();  
  
    ClearDataBreakpoint ( x );  
  
    if (breakpoint fired)  
        ReportDataRace();  
}
```

# Sampling Memory Accesses

- Need to execute the sampling function at every access
  - Performance overhead
- Need a good sampling algorithm
  - Data races are 'rare' events

# Dynamic Sampling

1. Toss a biased coin at each memory access
2. Sample if coin turns up heads (with small probability)

# Dynamic Sampling

1. Toss a biased coin at each memory access
  2. Sample if coin turns up heads (with small probability)
- Hot (frequent) instructions are sampled overwhelmingly more than cold instructions
  - Bugs (data races) are likely to be present on cold instructions

# Code Breakpoint Sampling

1. Pick a random instruction X
2. Set a code breakpoint at X
  - Overwrite the first byte of X with a '0xcc'
3. Sample X when the breakpoint fires and goto Step 1

# Sampling Using Code Breakpoints

- Samples instructions independent of their execution frequency
  - Hot and cold instructions are sampled uniformly



# Sampling Using Code Breakpoints

- Samples instructions independent of their execution frequency
  - Hot and code instructions are sampled uniformly

```
repeat {  
    t = fair_coin_toss();  
    while( t != unfair_coin_toss() );  
    print( t );  
}
```

- Sampling distribution is determined by you (fair\_coin\_toss)
- Sampling rate is determined by the program (unfair\_coin\_toss)

# Sampling Using Code Breakpoints

- Samples instructions independent of their execution frequency
  - Hot and code instructions are sampled uniformly

```
repeat {  
    t = fair_coin_toss();  
    while( t != unfair_coin_toss() );  
    print( t );  
}
```

Set a breakpoint  
at location X

Run the program  
till it executes X

Sample X

- Sampling distribution is determined by you (fair\_coin\_toss)
- Sampling rate is determined by the program (unfair\_coin\_toss)

# The DataCollider Algorithm

```
Init:
    SetRandomCodeBreakpoints( n );

AtCodeBreakpoint( x ) {
    SetDataBreakpoint ( x );
    Delay();
    if (breakpoint fired)
        ReportDataRace();
    ClearDataBreakPoint ( x );
    SetRandomCodeBreakpoints( 1 );
}
```

# DataCollider Summary

- Tunable runtime overhead
  - Tune code breakpoint rate so that we get  $k$  samples per second
- Finds lots of data races in practice
- The fastest implementation
  - Someone hacked it up during the talk
- Pruning intended data races is a problem
- Sampling algorithm improvements?