

# LOCKING PROTOCOLS & SOFTWARE TRANSACTIONAL MEMORY

---

# Why Study (Software) Transactions

# Why Study (Software) Transactions

1. Transactions are good abstractions

# Why Study (Software) Transactions

1. Transactions are good abstractions
2. Design choices and issues can guide your concurrent data structures

# Coarse/Fine Grained Locking

- Locking discipline enforce protection
  - Each object is protected by some lock
  - One lock protects a set of objects
- Lock leveling for deadlock avoidance
  - Assign a partial ordering to locks,  $\text{Level}(X)$
  - After acquiring a lock  $X$ , the transaction is only allowed to acquire locks  $Y$  with  $\text{Level}(Y) > \text{Level}(X)$
- Enforce locking discipline and lock leveling using static analysis tools

## Two Phase Locking (Revisited)

1. Transactions acquire a lock on an object before accessing the object
  2. Transactions cannot acquire a lock once the transaction has released some lock
- Works with shared/exclusive locks
  - Deadlock avoidance

# Use Hardware Primitives

```
Lock (L) ;  
x = x + 1 ;  
Unlock (L) ;
```



```
InterlockedIncrement (&x) ;
```

- Interlocked operations are implemented by low-level locking by the hardware

```
asm{  
    lock inc [x] ;  
}
```

## Using Compare And Exchange

- If you need to protect more than one logical variables
- But the variables can be compressed into a 32/64 bits

```
Lock(L) ; // protects x and y
```

```
x = f(x, y) ;
```

```
y = g(x, y) ;
```

```
Unlock(L) ;
```



# Using Compare and Exchange

- Compress x and y into a single 'state' variable

```
retry:
s = state; // compress state into 64 bits

x = s.x; y = s.y;
x = f(x,y); y = g(x,y)
t.x = x; t.y = y;

if(! CompareAndExchange(&state, t, s))
    goto retry;
```

# Progress Guarantees

- Wait freedom
  - Every thread makes progress in the presence of conflicts
- Lock freedom
  - Some thread makes progress in the presence of conflicts
- Obstruction freedom
  - A thread makes progress when all other threads are suspended

# What is the progress guarantee?

```
Lock (L) ;
```

```
x = f(x, y) ;
```

```
y = g(x, y) ;
```

```
Unlock (L) ;
```

```
Lock (L) ;
```

```
x = f(x, y) ;
```

```
y = g(x, y) ;
```

```
Unlock (L) ;
```

# What is the progress guarantee?

```
retry:
s = state;
x = s.x; y = s.y;
x = f(x,y);
y = g(x,y)
t.x = x; t.y = y;

if(!CAS(&state,t,s))
    goto retry;
```

```
retry:
s = state;
x = s.x; y = s.y;
x = f(x,y);
y = g(x,y)
t.x = x; t.y = y;

if(!CAS(&state,t,s))
    goto retry;
```

## Sophisticated Fine Grained Locking Protocols

- Locking protocols for guaranteeing linearizability
  - Even when not following the 2-phase locking discipline
- Allows early release of locks for more concurrency
- Proof of correctness is complicated
  - Never roll out your own

# Tree Locking


- The objects form a tree
  - Transactions acquire a lock on an object before accessing the object
    - Henceforth not state this constraint explicitly
1. A transaction may begin executing by locking any object
  2. Can acquire a lock on X only if already holding the lock on parent of X
  3. Can acquire a lock on X at most once
  4. (Can otherwise release locks in arbitrary order)

# “Hand Over Hand” Locking

```
void IncrementAll(List list, int k)
{
    Node c = list;


    while( c != null ) {
        c.data += k;

        c = c.next ;
    }
}
```



# “Hand Over Hand” Locking

```
void IncrementAll(List list, int k)
{
    Node c = list;
    Lock(c);
    while( c != null ) {
        c.data += k;
        if(c.next != null)
            Lock(c.next);
        c = c.next ;
        Unlock(c);
    }
}
```





# Proof of Tree Locking


- If two transactions conflict, one of the two acquires all shared locks before the other
- This ordering enforces the serialization order
  - => Serializability
- Is this protocol linearizable?

# DAG Locking

- The objects form a DAG
  1. A transaction may begin executing by locking any object
  2. Can acquire a lock on X only if
    1. already holding the lock on some parent of X
    2. Has previously acquired locks on all parents of X
  3. Can acquire a lock on X at most once
  4. (Can otherwise release locks in arbitrary order)

# Domination Locking


- Leverages the restricted semantics of software modules
  - thread can access n3 only after n1 & n2
- Allows early release
- Allows for dynamic modification to the underlying structure



## Two types of objects

- Distinguishes between **exposed** and **hidden** objects
- **Exposed** objects
  - “roots” of data structures
  - may be pointed by transaction arguments
- **Hidden** objects
  - may not be pointed by transaction arguments
  - may be reachable via exposed objects

***void insert(List l, int k) {...}***



# Domination Protocol Rules

1. A hidden object  $u$  can be acquired by  $t$  only if every path between an exposed object to  $u$  includes an object which is locked by  $t$
2. Perform leveled two phase locking on exposed objects

# Software Transactional Memory

- Different tradeoffs than database transactions
- Disk access  $\sim 10\text{ms}$  vs Memory access  $\sim 10\text{ns}$
- Durability (the D in ACID is not necessary)
- Databases can control all access to the data, while STMs should deal with non-transactional accesses

# Pessimism vs Optimism

- Pessimistic concurrency control
  - Expect conflict
  - Prevent conflict (by holding locks)
  - Need to avoid deadlocks
  - Useful when conflicts are frequent
- Optimistic concurrency control
  - Expect no conflicts
  - Rollback and retry
  - Need to avoid livelocks, wasted work on retry
  - Useful when conflicts are rare
- Combination possible

# What is the concurrency control in TL2?



## Version Management : How to perform writes

- “Eager” writes
  - Perform writes in place
  - Maintain undo-logs to revert writes on abort
- “Lazy” writes
  - Perform writes in a local log
  - Write to memory at commit time
  - Need to guarantee read-your-own-writes

# Dealing with inconsistent reads

```
atomic {  
    X = 1;  
  
    Y = 1;  
}
```

```
atomic {  
    if(X != Y)  
        while(1);  
}
```

# Obstruction Freedom:

## Is it acceptable to hold locks?

- Long running transactions can block progress ?
- A switched out thread can block progress ?
- A failed thread can block progress ?
- Priority inversion?

# Optimistic Concurrency For Reads

- Work loads tend to be read mostly
- Should allow efficient concurrent reads
- Look out for writes introduced in a read-only transaction

# TL2 Algorithm

```
read_ver = global_clock;
on_read(x):
    (x_ver, x_locked) = version(x);
    validate_read(x):
        if(x_locked || x_ver > read_ver) abort;
on_write(x):
    write (x,v) to log

forall w in write_set lock(w);
atomic{ write_ver = ++global_clock; }
forall r in read_set valid_read(r);
forall w in write_set{
    version(w) = write_ver;
    unlock(w);
}
```