

MAP REDUCE

Word Count Example

- We have a large file of words, one word to a line
- Count the number of times each distinct word appears in the file
- Sample application: analyze web server logs to find popular URLs

Word Count (2)

- Case 1: Entire file fits in memory

Word Count (2)

- Case 1: Entire file fits in memory
- Case 2: File too large for mem, but all <word, count> pairs fit in mem

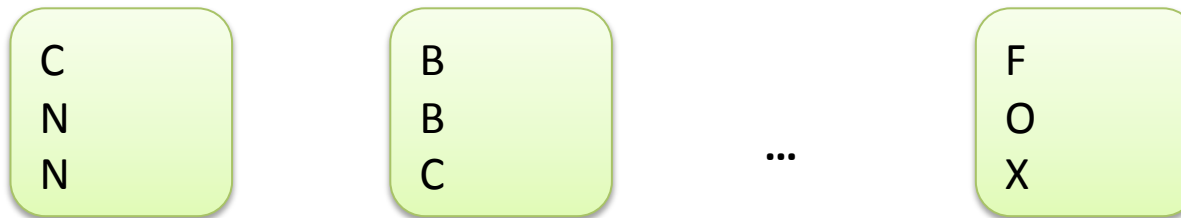
Word Count (2)

- Case 1: Entire file fits in memory
- Case 2: File too large for mem, but all <word, count> pairs fit in mem
- Case 3: File on disk, too many distinct words to fit in memory
 - **`sort datafile | uniq -c`**


Word Count (3)

- A large corpus of documents, sharded across many disks in many machines
- Machines, disks, networks can fail
- Motivation for Google's MapReduce

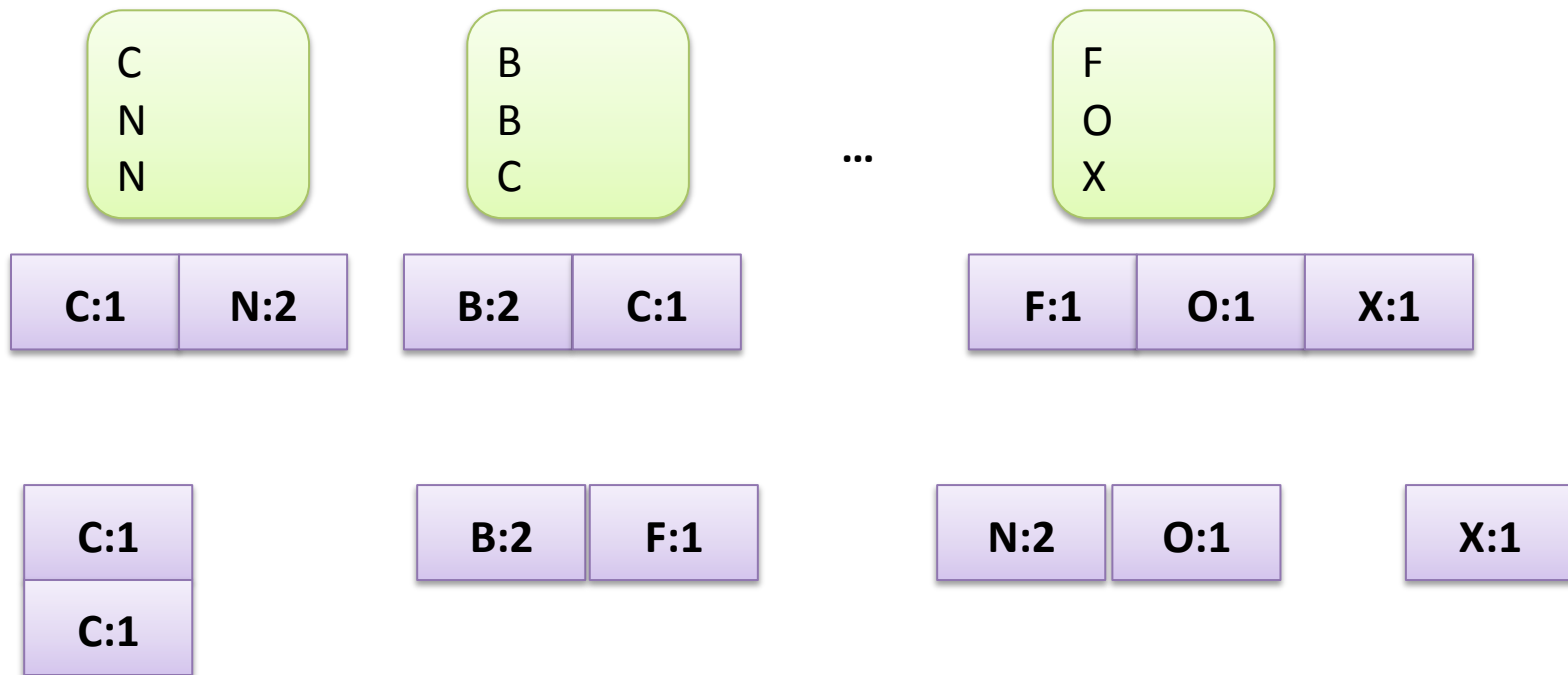
Input: A set of files



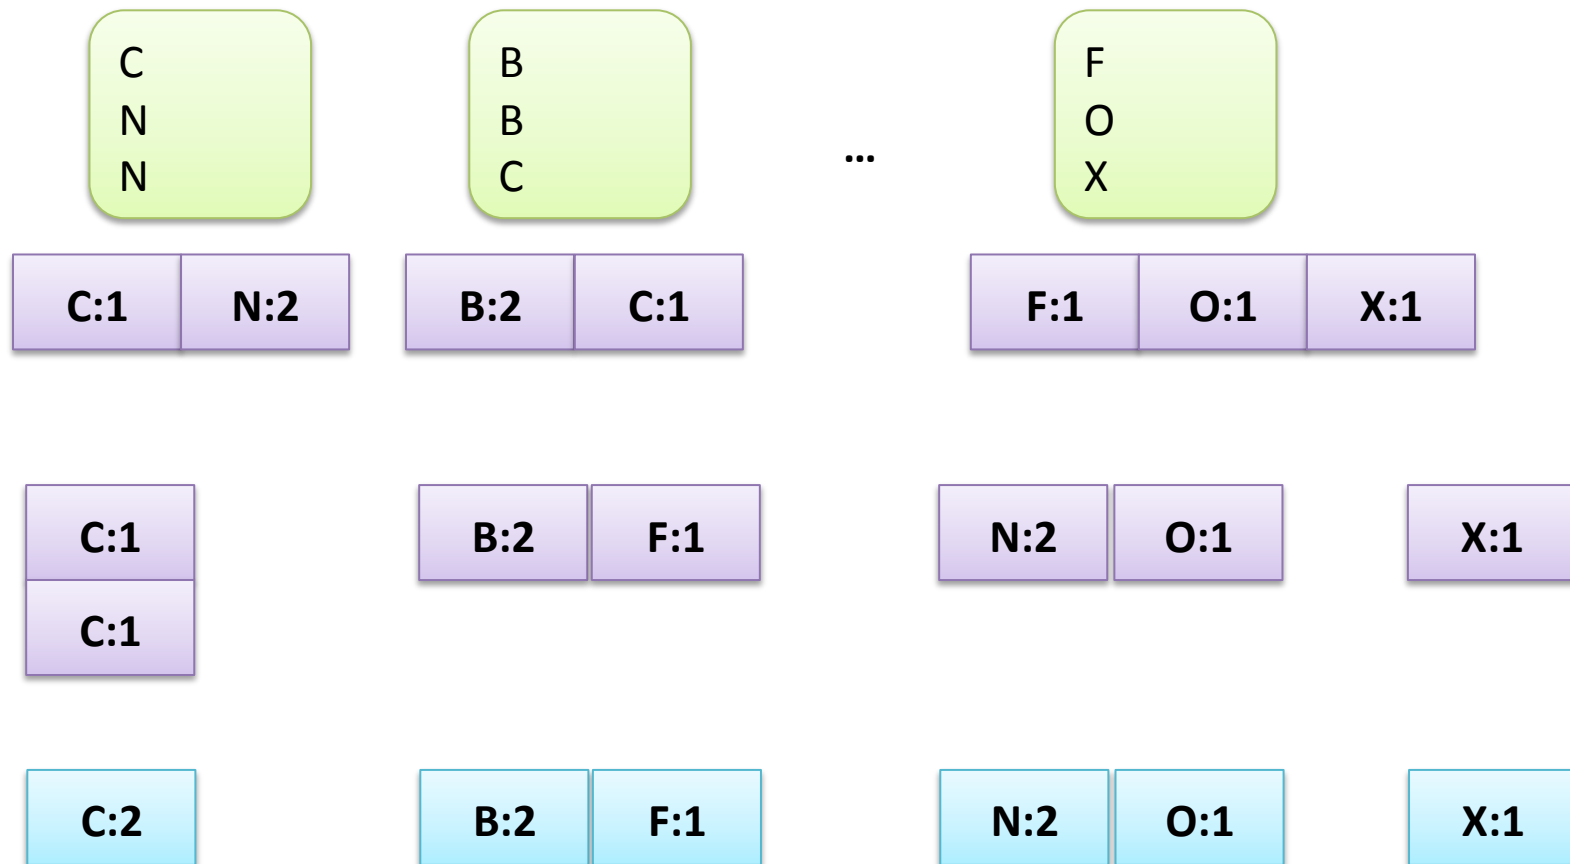
Map: Generate Word Count Per File



Partition (Optional)



Reduce



MapReduce

- Input: a set of key/value pairs
- User supplies two functions:
 - $\text{map}(k,v) \rightarrow \text{list}(k1,v1)$
 - $\text{reduce}(k1, \text{list}(v1)) \rightarrow v2$
- $(k1,v1)$ is an intermediate key/value pair
- Output is the set of $(k1,v2)$ pairs

Word Count using MapReduce

```
map(key, value):
```

```
// key: document name; value: text of document
```

```
  for each word w in value:
```

```
    emit(w, 1)
```

```
reduce(key, values):
```

```
// key: a word; value: an iterator over counts
```

```
  result = 0
```

```
  for each count v in values:
```

```
    result += v
```


```
  emit(result)
```

Map and Reduce vs MapReduce

- The map and reduce operations in MapReduce are inspired by similar operations in functional programming

Map

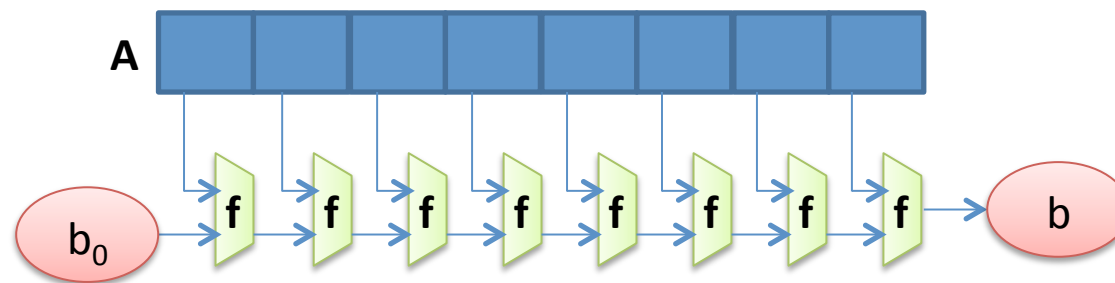
- Given a function $f : (A) \Rightarrow B$
- A collection $a: A[]$
- Generates a collection $b: B[]$, where $B[i] = f(A[i])$



- Parallel.For, Paralle.ForEach
 - Where each loop iteration is independent

Reduce

- Given a function $f: (A, B) \Rightarrow B$
- A collection $a: A[]$
- An initial value $b_0: B$
- Generate a final value $b: B$
 - Where $b = f(A[n-1], \dots f(A[1], f(A[0], b_0)))$



Relationship to SQL

- Implementing word count in SQL

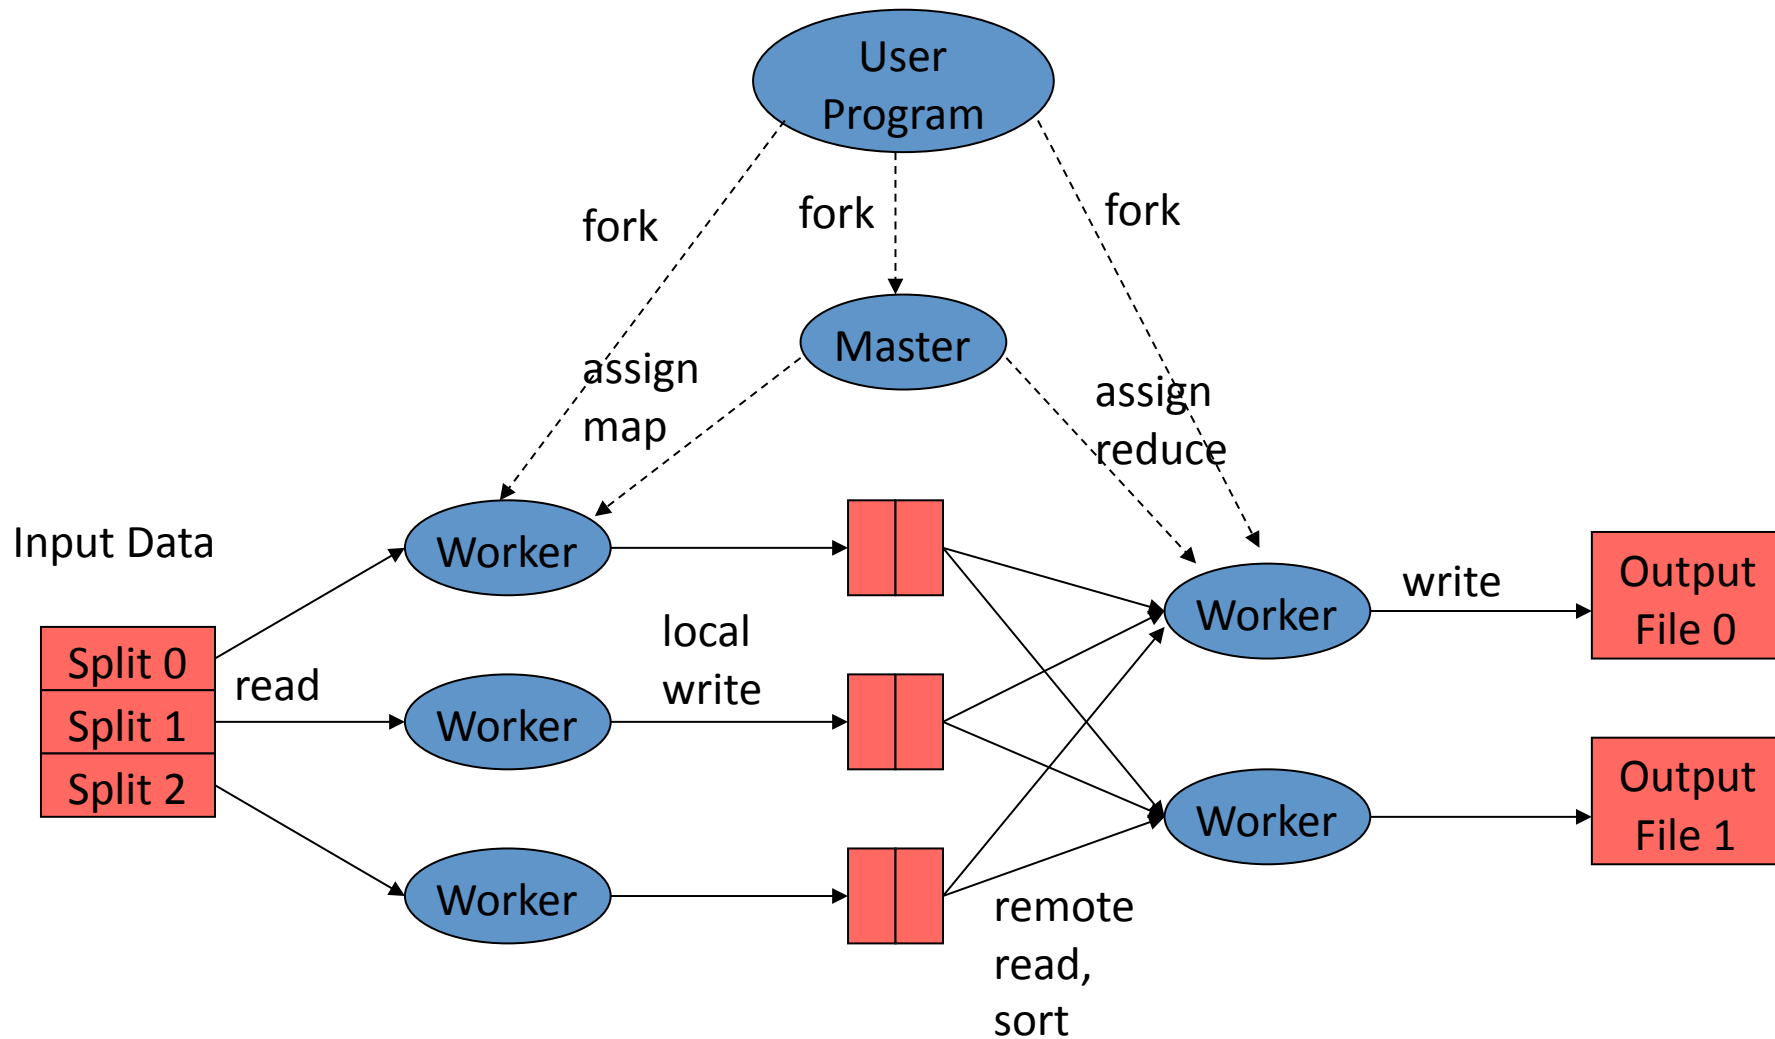
```
SELECT word Count(*) as wordCount  
FROM files  
GROUP BY word;
```

```
// where files is a (distributed)  
// relation <name, posn, word>
```


Signs of a Good Abstraction

- Hides important details
 - But not too much
- Simple for lay programmers to use
- Not necessarily general
 - But not very restricted
 - Can be application/domain specific
- Allows efficient implementations
 - Automatic optimizations
 - Manual optimizations (by experts)

Distributed Execution Overview



Data flow

- Input, final output are stored on a distributed file system
 - Scheduler tries to schedule map tasks “close” to physical storage location of input data
- Intermediate results are stored on local FS of map and reduce workers
- Output is often input to another map reduce task

Coordination

- Master data structures
 - Task status: (idle, in-progress, completed)
 - Idle tasks get scheduled as workers become available
 - When a map task completes, it sends the master the location and sizes of its R intermediate files, one for each reducer
 - Master pushes this info to reducers
- Master pings workers periodically to detect failures

Failures

- Map worker failure
 - Map tasks completed or in-progress at worker are reset to idle
 - Reduce workers are notified when task is rescheduled on another worker
- Reduce worker failure
 - Only in-progress tasks are reset to idle
- Master failure
 - MapReduce task is aborted and client is notified

How many Map and Reduce jobs?

- M map tasks, R reduce tasks
- Rule of thumb:
 - Make M and R much larger than the number of nodes in cluster
 - One DFS chunk per map is common
 - Improves dynamic load balancing and speeds recovery from worker failure
- Usually R is smaller than M, because output is spread across R files

Combiners

- Often a map task will produce many pairs of the form $(k, v_1), (k, v_2), \dots$ for the same key k
 - E.g., popular words in Word Count
- Can save network time by pre-aggregating at mapper
 - $\text{combine}(k_1, \text{list}(v_1)) \rightarrow v_2$
 - Usually same as reduce function
- Works only if reduce function is commutative and associative

Partition Function

- Inputs to map tasks are created by contiguous splits of input file
- For reduce, we need to ensure that records with the same intermediate key end up at the same worker
- System uses a default partition function e.g., $\text{hash}(\text{key}) \bmod R$
- Sometimes useful to override
 - E.g., $\text{hash}(\text{hostname}(\text{URL})) \bmod R$ ensures URLs from a host end up in the same output file

Avoiding Stragglers

- A slow running task (straggler) can prolong overall execution
 - Overloaded machines
 - Slow disk
- Kill stragglers
- Fork redundant tasks and take the first

Example: Sorting

Example: Database Join

Can Mappers Push instead of Reducers Pulling Data ?