

Patterns for Programming Shared Memory


Keep Abstraction Level High

- Temptation: ad-hoc parallelization
 - Add tasks or parallel loops all over the code
 - Discover data races/deadlocks, fix them one-by-one
- Problem:
 - Complexity quickly adds up
 - Easy to get cornered by deadlocks, atomicity violations, and data races
 - And these bugs often are hard to expose
- Alternatives?
 - Use well-understood, simple high-level patterns

Book Recommendation:

Parallel Programming with Microsoft .NET

- Covers Common Patterns Systematically
- Free download at <http://parallelpatterns.codeplex.com>



Sharing State Safely

- We discuss two types of patterns in this unit
 - Not an exhaustive list, read book for more

Architectural Patterns

Localize shared state

Producer-Consumer

Pipeline

Worklist

Replication Patterns

Make copies of shared state

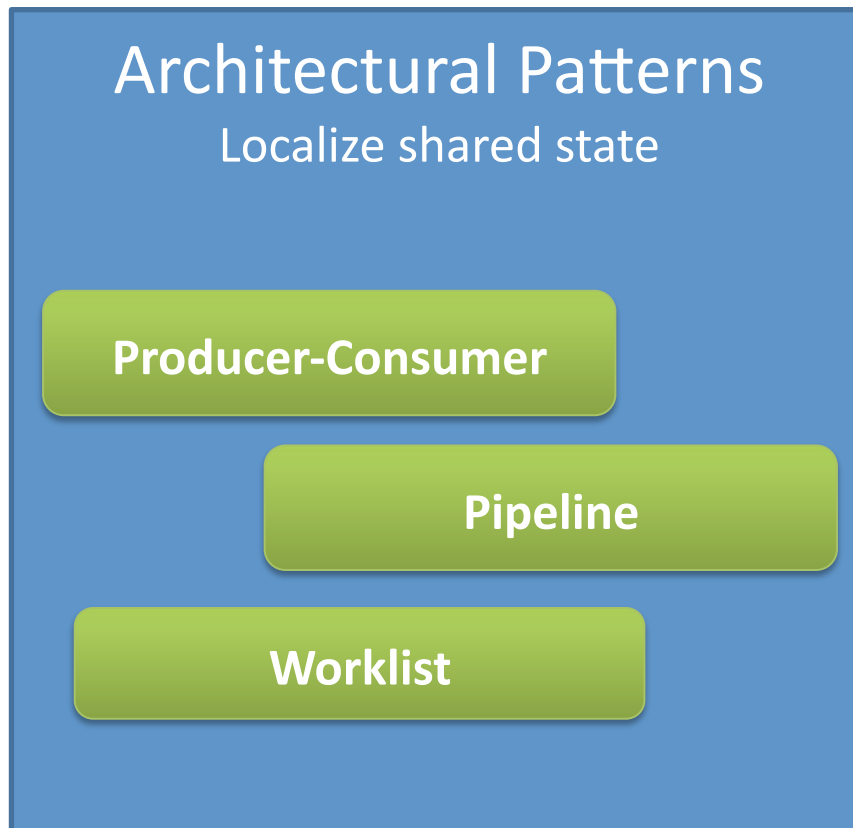
Immutable Data

Double Buffering

Part I

ARCHITECTURAL PATTERNS


Architectural Patterns



- Visualize “factory”: each worker works on a separate item
- Items travel between workers by means of buffers

Example: Producer-Consumer

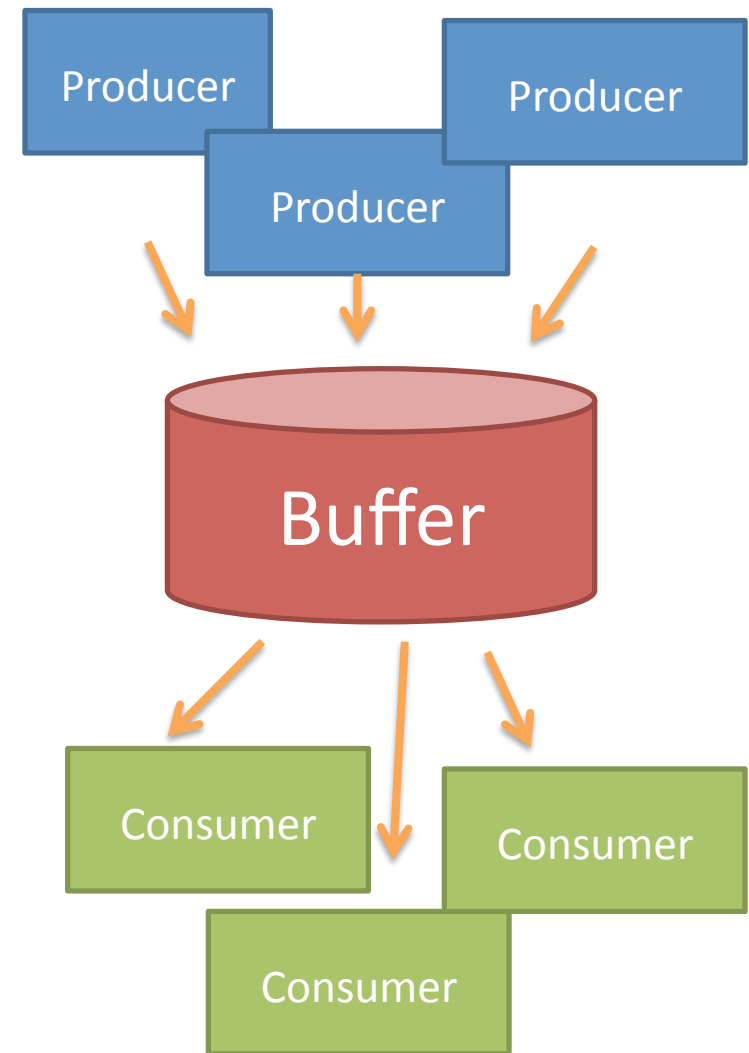
- One or more producers add items to the buffer
- One or more consumers remove items from the buffer



Example: Producer-Consumer

No data races!

1. Item is local to Producer before insertion into buffer
2. Item is local to Consumer after removal from buffer
3. What about buffer?
 - Buffer is thread-safe
 - Blocks when full/empty
 - Not trivial to implement
 - More about this in Unit 3
 - For now, use `BlockingCollection<T>`



Producer-Consumer Code

```
int buffersize = 1024;
int num_producers = 4;
int num_consumers = 4;

// create bounded buffer
var buffer = new BlockingCollection<Item>(buffersize);

// start consumers
for (int i = 0; i < num_consumers; i++)
    Task.Factory.StartNew(() => new Consumer(buffer).Run());

// start producers
for (int i = 0; i < num_producers; i++)
    Task.Factory.StartNew(() => new Producer(buffer).Run());
```

Producer Code

```
class Producer
{
    public Producer(BlockingCollection<Item>
buffer)
    {
        this.buffer = buffer;
    }
    public void Run()
    {
        while (!Done())
        {
            Item i = Produce();
            buffer.Add(i);
        }
    }
    private BlockingCollection<Item> buffer;
    private Item Produce() { ... }
    private bool Done() { ... }
}
```

Consumer Code

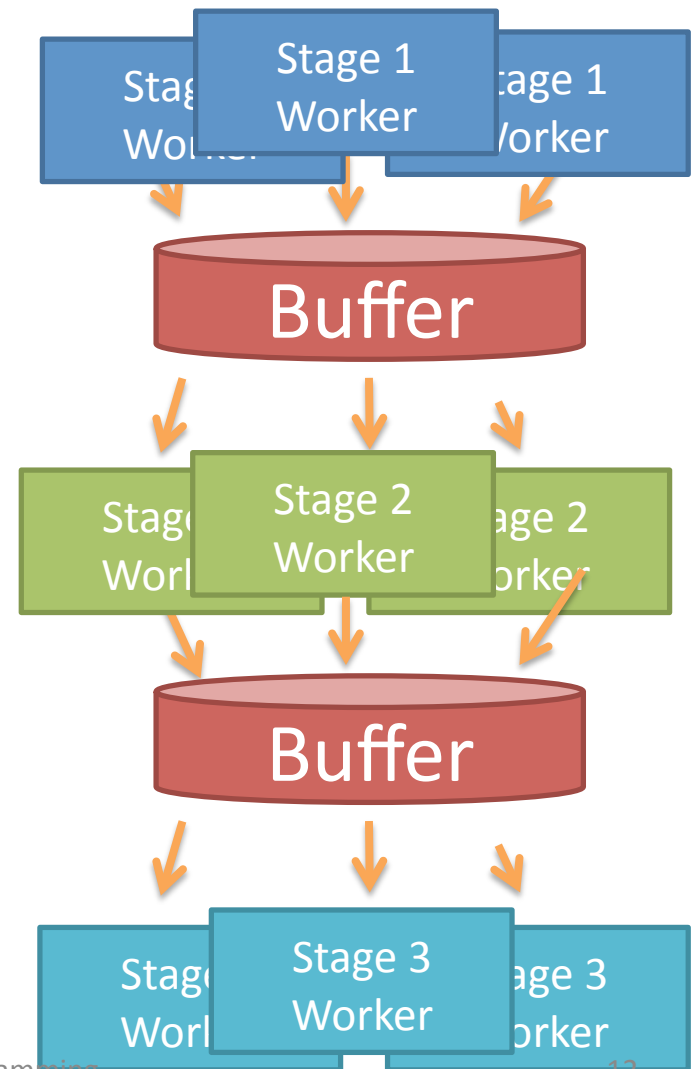
```
class Consumer
{
    public Consumer(BlockingCollection<Item> buffer)
    {
        this.buffer = buffer;
    }
    private BlockingCollection<Item> buffer;

    public void Run()
    {
        foreach (Item i in buffer.GetConsumingEnumerable())
            Consume(i);
    }
    private void Consume(Item item)
    {
        ...
    }
}
```

6/22/2010

Pipeline Pattern

- Generalization of Producer-Consumer
 - One or more workers per stage
 - First stage = Producer
 - Last stage = Consumer
 - Middle stages consume and produce
 - **No Data Races:** Data is local to workers



Pipeline Code

```
// create bounded buffers
var buffer_1_2 = new BlockingCollection<Item>(buffersize);
var buffer_2_3 = new BlockingCollection<Item>(buffersize);

// start workers
for (int i = 0; i < num_workers_stage1; i++)
    Task.Factory.StartNew(() =>
        new Stage1Worker(buffer_1_2).Run());
for (int i = 0; i < num_workers_stage2; i++)
    Task.Factory.StartNew(() =>
        new Stage2Worker(buffer_1_2, buffer_2_3).Run());
for (int i = 0; i < num_workers_stage3; i++)
    Task.Factory.StartNew(() =>
        new Stage3Worker(buffer_2_3).Run());
```


Stage Worker Code

```
class Stage2Worker {
    public Stage2Worker(BlockingCollection<Item> buffer_in,
                        BlockingCollection<Item> buffer_out) {
        this.buffer_in = buffer_in;
        this.buffer_out = buffer_out;
    }

    public void Run() {
        try {
            foreach (Item i in buffer_in.GetConsumingEnumerable()) {
                Item j = Process(i);
                buffer_out.Add(j);
            }
        }
        finally { buffer_out.CompleteAdding(); }
    }
    . . .
}
```

Worklist Pattern

- Worklist contains items to process
 - Workers grab one item at a time
 - Workers may add items back to worklist
 - No data races: items are local to workers



Worklist Code (1/2)

```
public void Run()
{
    int num_workers = 4;

    // create worklist, filled with initial work
    worklist = new BlockingCollection<Item>(
        new ConcurrentQueue<Item>(GetInitialWork()));

    cts = new CancellationTokenSource();
    itemcount = worklist.Count();

    for (int i = 0; i < num_workers; i++)
        Task.Factory.StartNew(RunWorker);
}

IEnumerable<Item> GetInitialWork() { ... }
```

```
BlockingCollection<Item> worklist;
CancellationTokenSource cts;
int itemcount;
```


Worklist Code (2/2)

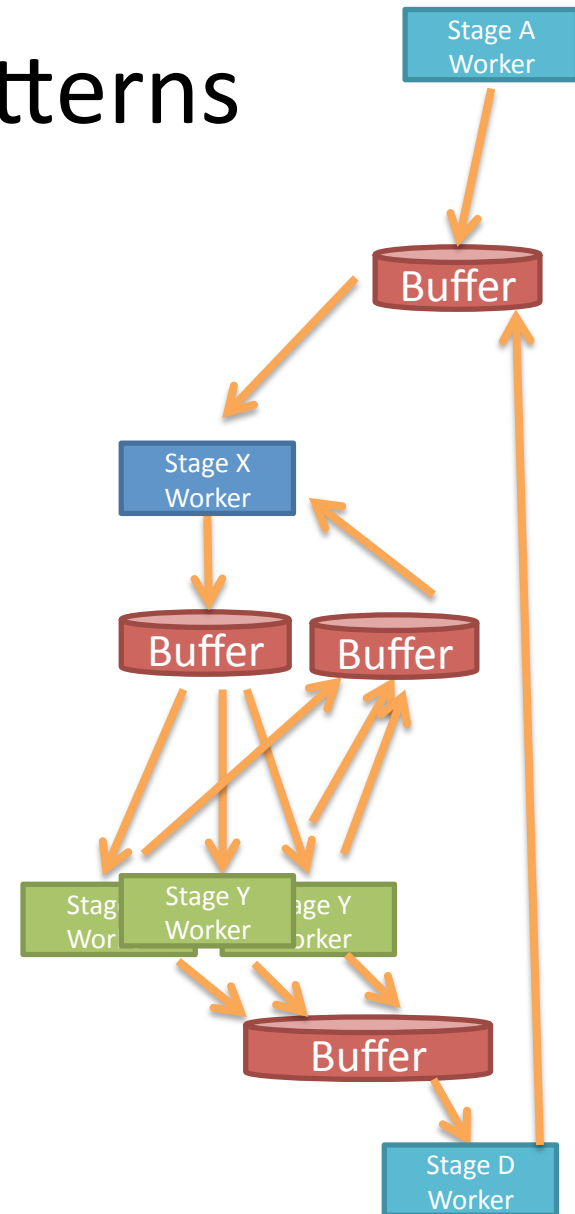
```
public void RunWorker() {
    try {
        do {
            Item i = worklist.Take(cts.Token);
            // blocks until item available or cancelled
            Process(i);
            // exit loop if no more items left
        } while (Interlocked.Decrement(ref itemcount) > 0);
    } finally {
        if (!cts.IsCancellationRequested)
            cts.Cancel();
    }
}

public void Process(Item i) { . . . }
// may call AddWork() to add more work to worklist
public void AddWork(Item item) {
    Interlocked.Increment(ref itemcount);
    worklist.Add(item);
}
```

```
BlockingCollection<Item> worklist;
CancellationTokenSource cts;
int itemcount;
```

Application Architecture May Combine Patterns

- “Stages” arranged in a graph and may have
 - Single worker (only inter-stage parallelism)
 - Multiple workers (also intra-stage parallelism)
- No data races: items local to worker
- Buffers may have various ordering characteristics
 - ConcurrentQueue = FIFO
 - ConcurrentStack = LIFO
 - ConcurrentBag = unordered



Provisioning

- How to choose number of workers?
 - Single worker
 - If ordering of items matter
 - Workers are CPU-bound
 - Can not utilize more than one per processor
`num_workers = Environment.ProcessorCount`
 - Workers block in synchronous IO
 - Can potentially utilize many more
- How to choose buffer size?
 - Too small -> forces too much context switching
 - Too big -> uses inappropriate amount of memory

Manual Provisioning

```
BlockingCollection<Item> buffer;
class Item { ... }
void Consume(Item i) { ... }

public void FixedConsumers() {
    int num_workers = Environment.ProcessorCount;
    for (int i = 0; i < num_workers; i++)
        Task.Factory.StartNew(() => {
            foreach (Item item in buffer.GetConsumingEnumerable())
                Consume(item);
        });
}
```

Automatic Provisioning

```
BlockingCollection<Item> buffer;  
class Item { ... }  
void Consume(Item i) { ... }
```

```
public void FlexibleConsumers()  
{  
    Parallel.ForEach(  
        new BlockingCollectionPartitioner<Item>(buffer),  
        Consume);  
}
```

```
private class BlockingCollectionPartitioner<T> : Partitioner<T>  
{  
    // see http://blogs.msdn.com/b/pfxteam/archive/2010/04/06/9990420.aspx  
}
```

Part II

REPLICATION PATTERNS

Replication Patterns

- Idea: Many workers can work on the “same” item if we make copies

Replication Patterns

Make copies of shared state

Immutable Data

Double Buffering

Immutability

- Remember: concurrent reads do not conflict
- Idea: never write to shared data
 - All shared data is immutable (read only)
 - To modify data, must make a fresh copy first
- Example: strings
 - Strings are immutable!
 - Never need to worry about concurrent access of strings.

Build your own immutable objects

- Mutable list: changes original

```
interface IMutableSimpleList<T>
{
    int Count { get; } // read only property
    T this[int index] { get; } // read only indexer
    void Add(T item); // mutator method
}
```

- Immutable list: never changes original, copies on mutation

```
interface IImmutableSimpleList<T>
{
    int Count { get; } // read only property
    T this[int index] { get; } // read only indexer
    ImmutableSimplelist<T> Add(T item); // read only method
}
```

Copy on Write

```
class SimpleImmutableList<T> : IImmutableSimpleList<T> {  
    private T[] _items = new T[0];  
  
    public int Count { get { return _items.Length; } }  
    public T this[int index] { get { return _items  
[index]; } }  
  
    public IImmutableSimpleList<T> Add(T item)  
    {  
        T[] items = new T[_items.Length + 1];  
        Array.Copy(_items, items, _items.Length);  
        items[items.Length - 1] = item;  
  
        return new SimpleImmutableList<T>(items);  
    }  
}
```

Example: Game

```
class GameObject {
    public void Draw()
    { // ... reads state
    }
    public void Update()
    { // ... modifies state
    }
}

GameObject[] objects
= new GameObject[numobjects];

void DrawAllCells()
{
    foreach(GameObject g in objects)
        g.Draw();
}

void UpdateAllCells()
{
    foreach(GameObject g in objects)
        g.Update();
}
```

```
// GOAL: draw and update in parallel
void GameLoop()
{
    while (!Done()) {
        DrawAllCells();
        UpdateAllCells();
    }
}
```

```
// incorrect parallelization:
// data race between draw and update
void ParallelGameLoop()
{
    while (!Done())
        Parallel.Invoke(DrawAllCells,
                        UpdateAllCells)
}
```

Can we use immutability to fix this?

Example: Game

```
class GameObject {  
    public void Draw()  
    { // ... reads state  
    }  
    public GameObject ImmutableUpdate()  
    { // ... return updated copy  
    }  
}
```

```
GameObject[] objects  
    = new GameObject[numobjects];
```

```
void DrawAllCells()  
{  
    foreach(GameObject g in objects)  
        g.Draw();  
}
```

```
GameObject[] UpdateAllCells()  
{  
    var newobjects = new GameObject[numobjects];  
    for (int i = 0; i < numobjects; i++)  
        newobjects[i] = objects[i].ImmutableUpdate();  
    return newObjects;  
}
```

```
// correct parallelization of loop  
void ParallelGameLoop()  
{  
    while (!Done())  
    {  
        GameObject[] newarray = null;  
        Parallel.Invoke(() => DrawAllCells(),  
            () => newarray = UpdateAllCells());  
        objects = newarray;  
    }  
}
```

Correct parallelization.

Can we do this without
allocating a new array
every iteration?

Trick: Double Buffering

```
class GameObject {
    public void Draw()
    { // ... reads state
    }
    public GameObject ImmutableUpdate()
    { // ...
    }
}

GameObject[,] objects
    = new GameObject[2, numobjects];

void DrawAllCells(int bufferIdx)
{
    for (int i = 0; i < numobjects; i++)
        objects[bufferIdx, i].Draw();
}

void UpdateAllCells(int bufferIdx)
{
    for (int i = 0; i < numobjects; i++)
        objects[1 - bufferIdx, i] = objects[bufferIdx, i].ImmutableUpdate();
}
```

```
// correct parallelization of loop
void ParallelGameLoop()
{
    int bufferIdx = 0;
    while (!Done())
    {
        Parallel.Invoke(
            () => DrawAllCells(bufferIdx),
            () => UpdateAllCells(bufferIdx)
        );
        bufferIdx = 1 - bufferIdx;
    }
}
```