

Correctness of parallel programs

Shaz Qadeer
Research in Software Engineering

CSEP 506
Spring 2011

Why is correctness important?

- Software development is expensive
 - Testing and debugging significant part of the cost
- Testing and debugging of parallel and concurrent programs is even more difficult and expensive

The Heisenbug problem

- Sequential program: program execution fully determined by the input
- Parallel program: program execution may depend both on the input and the communication among the parallel tasks
- Communication dependencies are invariably not reproducible

Essence of reasoning about correctness

- Modeling
- Abstraction
- Specification
- Verification

What is a data race?

- An execution in which two tasks simultaneously access the same memory location

```
Parallel.For(0, filenames.Length, (int i) =>
{
    int len = filenames[i].Length;
    count[len]++;
});
```

Example

```
filenames.Length == 2
```

```
filenames[0].Length == filenames[1].Length == 1
```

```
Parallel.For(0, filenames.Length, (int i) =>  
    {  
        int len = filenames[i].Length;  
        count[len]++;  
    });
```

Task(0)

```
int len, t;
```

```
len = 1;
```

```
t = count[1];
```

```
t++;
```

```
count[1] = t;
```

Task(1)

```
int len, t;
```

```
len = 1;
```

```
t = count[1];
```

```
t++;
```

```
count[1] = t;
```

Two executions

E1

Task(0)

```
int len, t;
```

```
len = 1;
```

```
t = count[1];
```

```
t++;
```

```
count[1] = t;
```

Task(1)

```
int len, t;
```

```
len = 1;
```

```
t = count[1];
```

```
t++;
```

```
count[1] = t;
```

E2

Task(0)

```
int len, t;
```

```
len = 1;
```

```
t = count[1];
```

```
t++;
```

```
count[1] = t;
```

Task(1)

```
int len, t;
```

```
len = 1;
```

```
t = count[1];
```

```
t++;
```

```
count[1] = t;
```

Which execution has a data race?

Example

```
count.Length == 2
```

```
for (int i = 0; i < count.Length; i++) { count[i] = 0; }
```

```
Parallel.For(0, count.Length, (int i) => { count[i]++; });
```

```
for (int i = 0; i < count.Length; i++) { Console.WriteLine(count[0]); }
```


An execution

Task(0)

```
int t;
```

```
t = count[0];
```

```
t++;
```

```
count[0] = t;
```

Parent()

```
count[0] = 0;
```

```
count[1] = 0;
```

```
WriteLine(count[0]);
```

```
WriteLine(count[1]);
```

Task(1)

```
int t;
```

```
t = count[1];
```

```
t++;
```

```
count[1] = t;
```

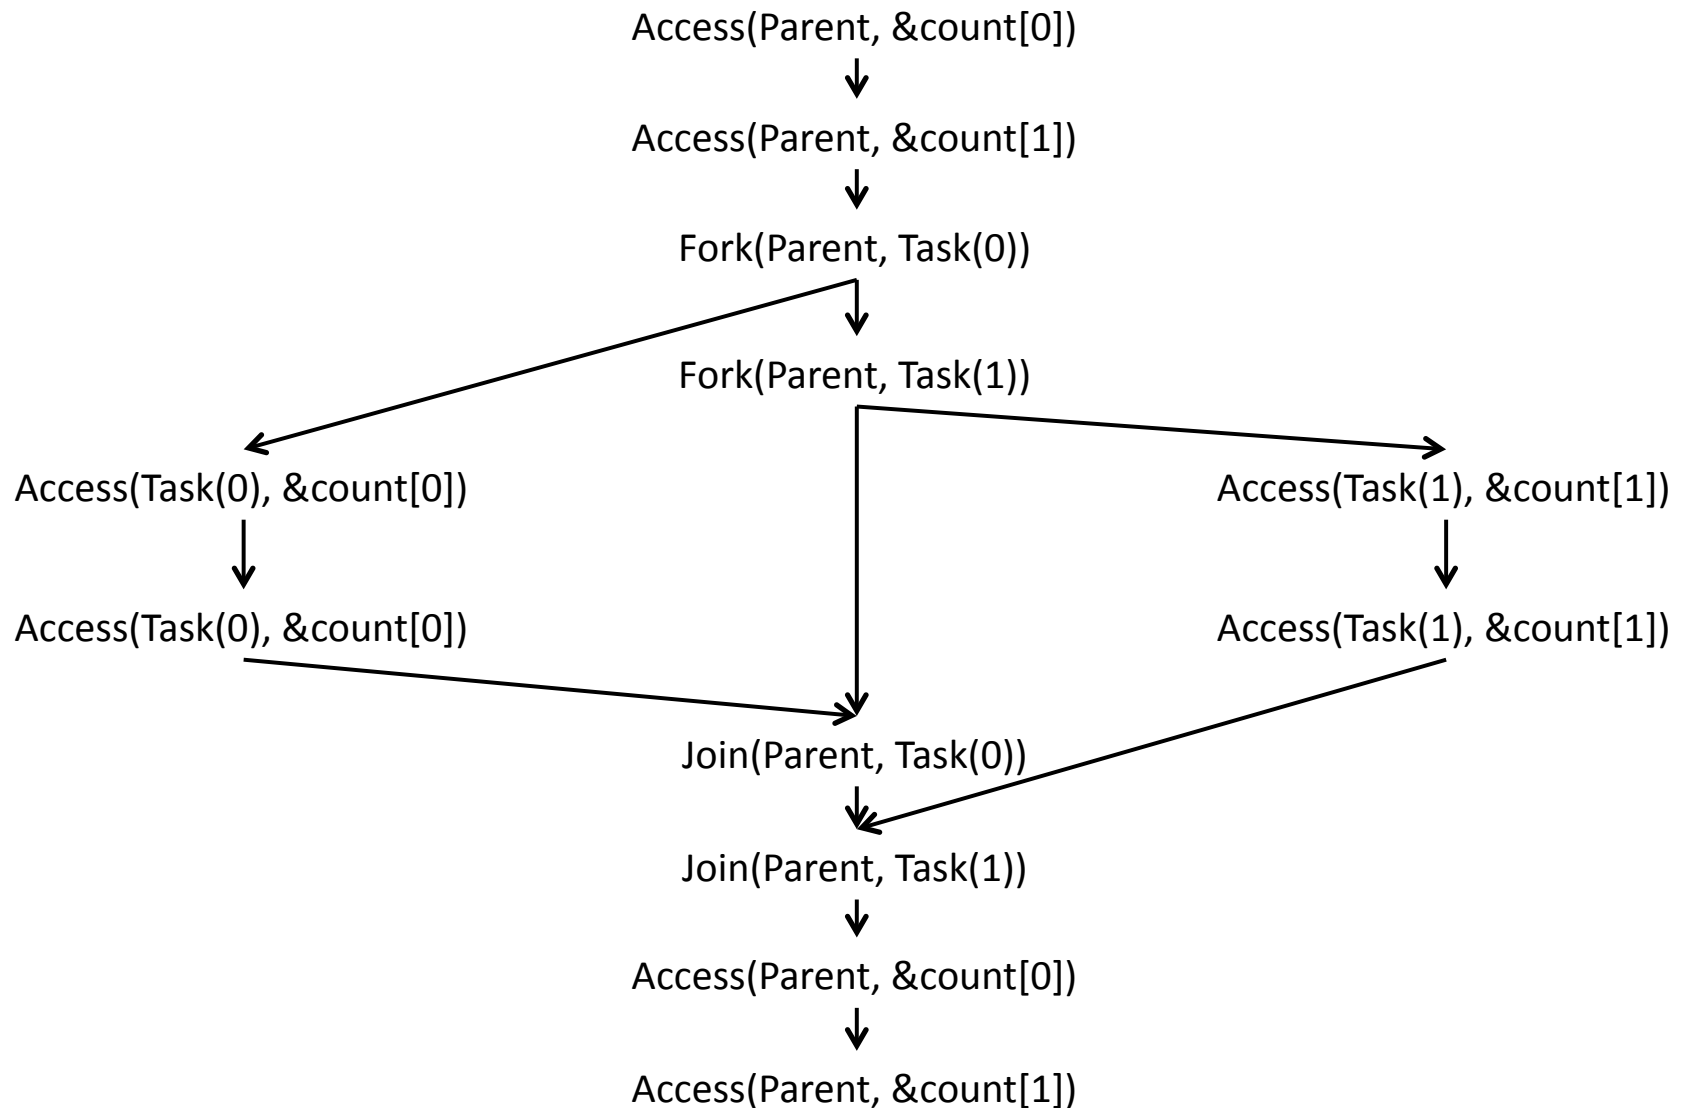
What is a parallel execution?

- Happens-before graph: directed acyclic graph over the set of events in the execution
- Three kinds of events
 - Access(t, a): task t accessed address a
 - Fork(t, u): task t created task u
 - Join(t, u): task t waited for task u
- Three kinds of edges
 - program order: edge from an event by a particular task to subsequent event by the same task
 - fork: edge from Fork(t, u) to first event performed by task u
 - join: edge from last event performed by task u to Join(t, u)

A note on modeling executions

- A real execution on a live system is complicated
 - instruction execution on the CPU
 - scheduling in the runtime
 - hardware events, e.g., cache coherence messages
- Focus on what is relevant to the specification
 - memory accesses because data-races are about conflicting memory accesses
 - fork and join operations because otherwise we cannot reason precisely

Example of happens-before graph



Definition of data race

- $e < f$ in an execution iff there is a path in the happens-before graph from e to f
 - e happens before f
- An execution has a data-race on address x iff there are different events e and f
 - both e and f access x
 - not $e < f$
 - not $f < e$
- An execution is race-free iff there is no data-race on any address x

Racy execution

```
Task(0)
int len, t;
len = 1;
t = count[1];
t++;
count[1] = t;
```

Access(Task(0),
&count[1])



Access(Task(0),
&count[1])

```
Task(1)
int len, t;
len = 1;
t = count[1];
t++;
count[1] = t;
```

Access(Task(1),
&count[1])



Access(Task(1),
&count[1])

```
Task(0)
int len, t;
len = 1;
t = count[1];
t++;
count[1] = t;
```

Access(Task(0),
&count[1])



Access(Task(0),
&count[1])


```
Task(1)
int len, t;
len = 1;
t = count[1];
t++;
count[1] = t;
```

Access(Task(1),
&count[1])



Access(Task(1),
&count[1])

Race-free execution



Vector-clock algorithm

- Vector clock: an array of integers indexed by task identifiers
- For each task t , maintain a vector clock $C(t)$
 - each clock in $C(t)$ initialized to 0
- For each address a , maintain a vector clock $X(a)$
 - each clock in $X(a)$ initialized to 0

Vector-clock operations

- Task t executes an event
 - increment $C(t)[t]$ by one
- Task t forks task u
 - initialize $C(u)$ to $C(t)$
- Task t joins with task u
 - update $C(t)$ to $\max(C(t), C(u))$
- Task t accesses address a
 - data race unless $X(a) < C(t)$
 - update $X(a)$ to $C(t)$

	C(Parent)	C(Task(0))	C(Task(1))	X(&count[0])	X(&count[0])
	[0, 0, 0]	[0, 0, 0]	[0, 0, 0]	[0, 0, 0]	[0, 0, 0]
Access(Parent, &count[0])	[1, 0, 0]			[1, 0, 0]	
Access(Parent, &count[1])	[2, 0, 0]				[2, 0, 0]
Fork(Parent, Task(0))	[3, 0, 0]	[3, 0, 0]			
Fork(Parent, Task(1))	[4, 0, 0]		[4, 0, 0]		
Access(Task(0), &count[0])		[3, 1, 0]		[3, 1, 0]	
Access(Task(1), &count[1])			[4, 0, 1]		[4, 0, 1]
Access(Task(0), &count[0])		[3, 2, 0]		[3, 2, 0]	
Access(Task(1), &count[1])			[4, 0, 2]		[4, 0, 2]
Join(Parent, Task(0))	[5, 2, 0]				
Join(Parent, Task(1))	[6, 2, 2]				
Access(Parent, &count[0])	[7, 2, 2]			[7, 2, 2]	
Access(Parent, &count[1])	[8, 2, 2]				[8, 2, 2]

	C(Task(0))	C(Task(1))	X(&count[0])	X(&count[0])
	[0, 0]	[0, 0]	[0, 0]	[0, 0]
Access(Task(0), &count[1])	[1, 0]			[1, 0]
Access(Task(1), &count[1])		[0, 1]		
Access(Task(0), &count[1])				
Access(Task(1), &count[1])				

Correctness of vector-clock algorithm

- For any execution and any two events $e@t$ followed by $f@u$ in the execution
 - $e@t < f@u$ iff $C[t]@e < C[u]@f$

Synchronizing using locks

```
Parallel.For(0, filenames.Length, (int i) =>
{
    int len = filenames[i].Length;
    lock (lock[len]) { count[len]++; }
});
```

```
filenames.Length == 2  
filenames[0].Length == filenames[1].Length == 1
```

```
count[1] = 0;  
Parallel.For(0, filenames.Length, (int i) =>  
    {  
        int len = filenames[i].Length;  
        lock (lock[len]) { count[len]++; }  
    });  
Console.WriteLine(count[1]);
```

Parent

```
count[1] = 0;  
Console.WriteLine(count[1]);
```

Task(0)

```
int len, t;  
  
len = 1;  
acquire(lock[1]);  
t = count[1];  
t++;  
count[1] = t;  
release(lock[1]);
```

Task(1)

```
int len, t;  
  
len = 1;  
acquire(lock[1]);  
t = count[1];  
t++;  
count[1] = t;  
release(lock[1]);
```

Execution I

Parent

```
count[1] = 0;
```

```
Console.WriteLine(count[1]);
```

Task(0)

```
int len, t;
```

```
len = 1;
```

```
acquire(lock[1]);
```

```
t = count[1];
```

```
t++;
```

```
count[1] = t;
```

```
release(lock[1]);
```

Task(1)

```
int len, t;
```

```
len = 1;
```

```
acquire(lock[1]);
```

```
t = count[1];
```

```
t++;
```

```
count[1] = t;
```

```
release(lock[1]);
```

Execution II

Parent

```
count[1] = 0;
```

```
Console.WriteLine(count[1]);
```

Task(0)

```
int len, t;
```

```
len = 1;
```

```
acquire(lock[1]);
```

```
t = count[1];
```

```
t++;
```

```
count[1] = t;
```

```
release(lock[1]);
```

Task(1)

```
int len, t;
```

```
len = 1;
```

```
acquire(lock[1]);
```

```
t = count[1];
```

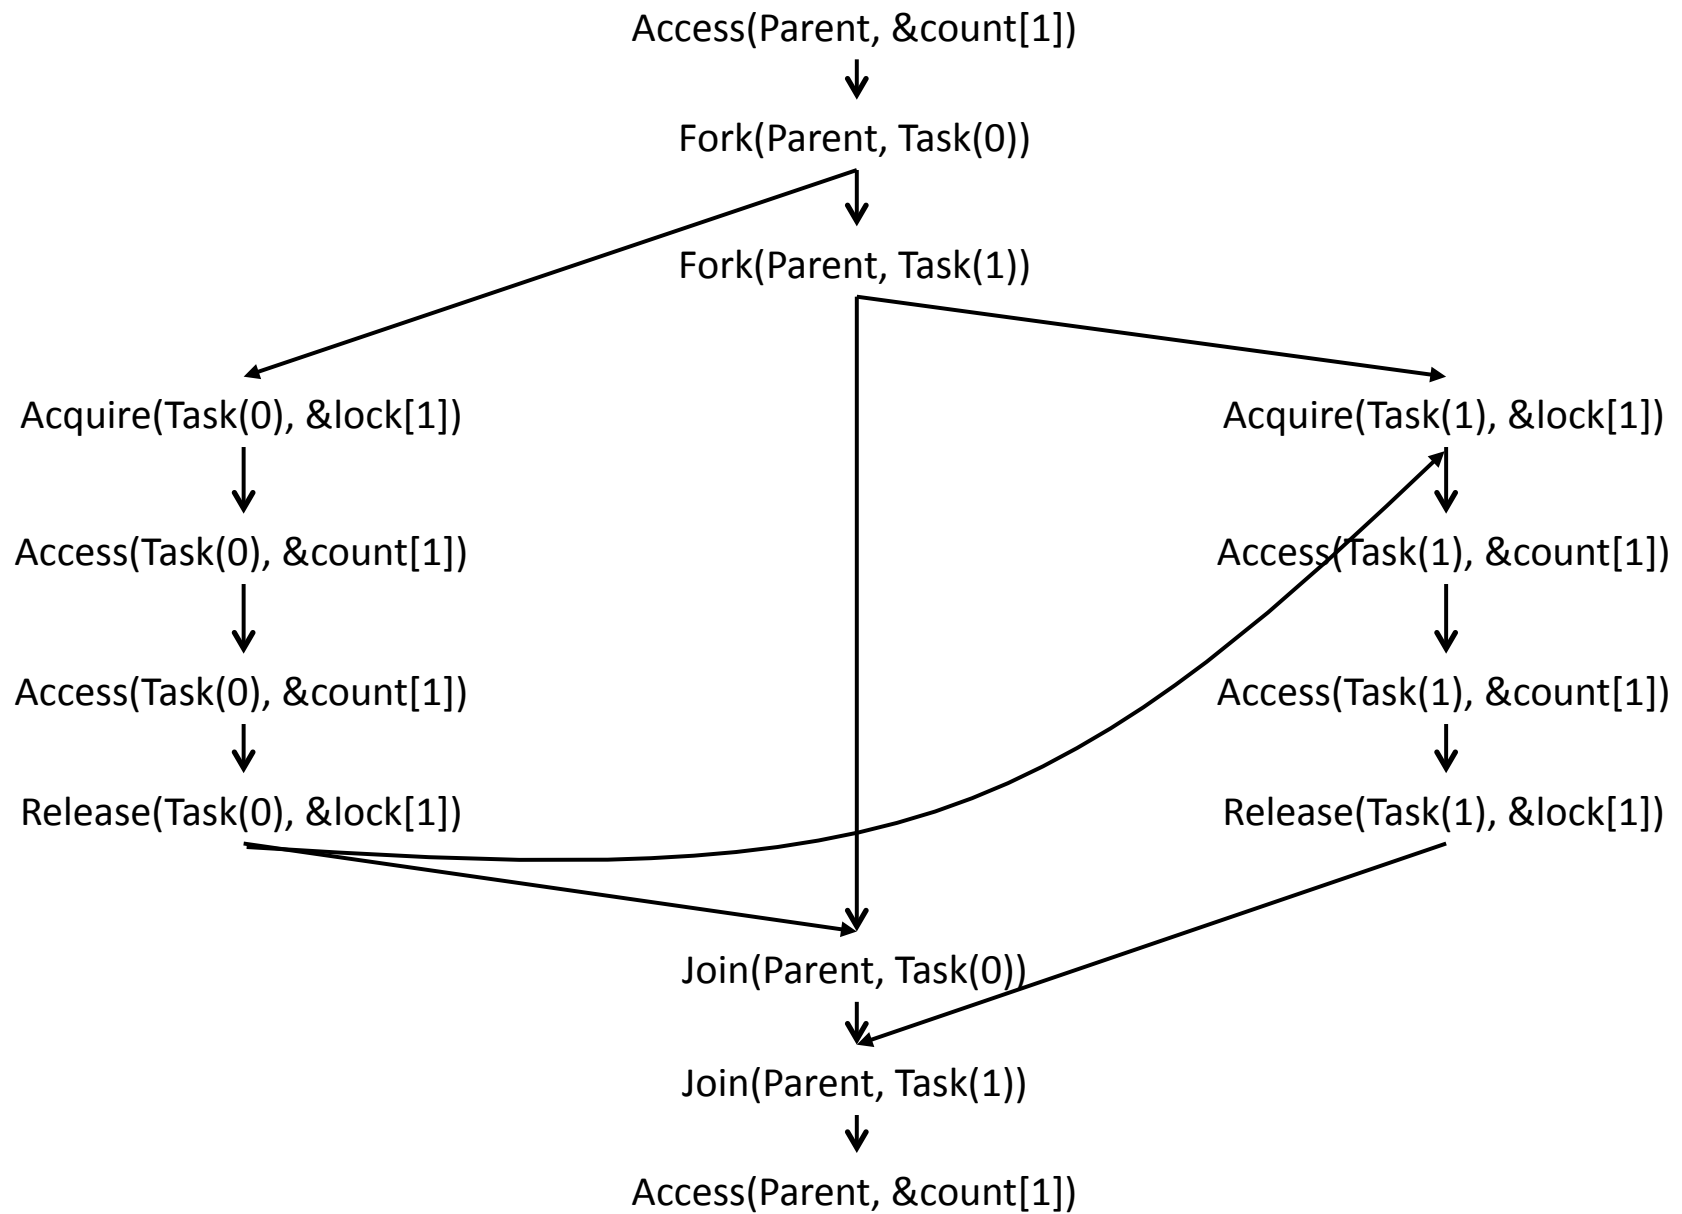
```
t++;
```

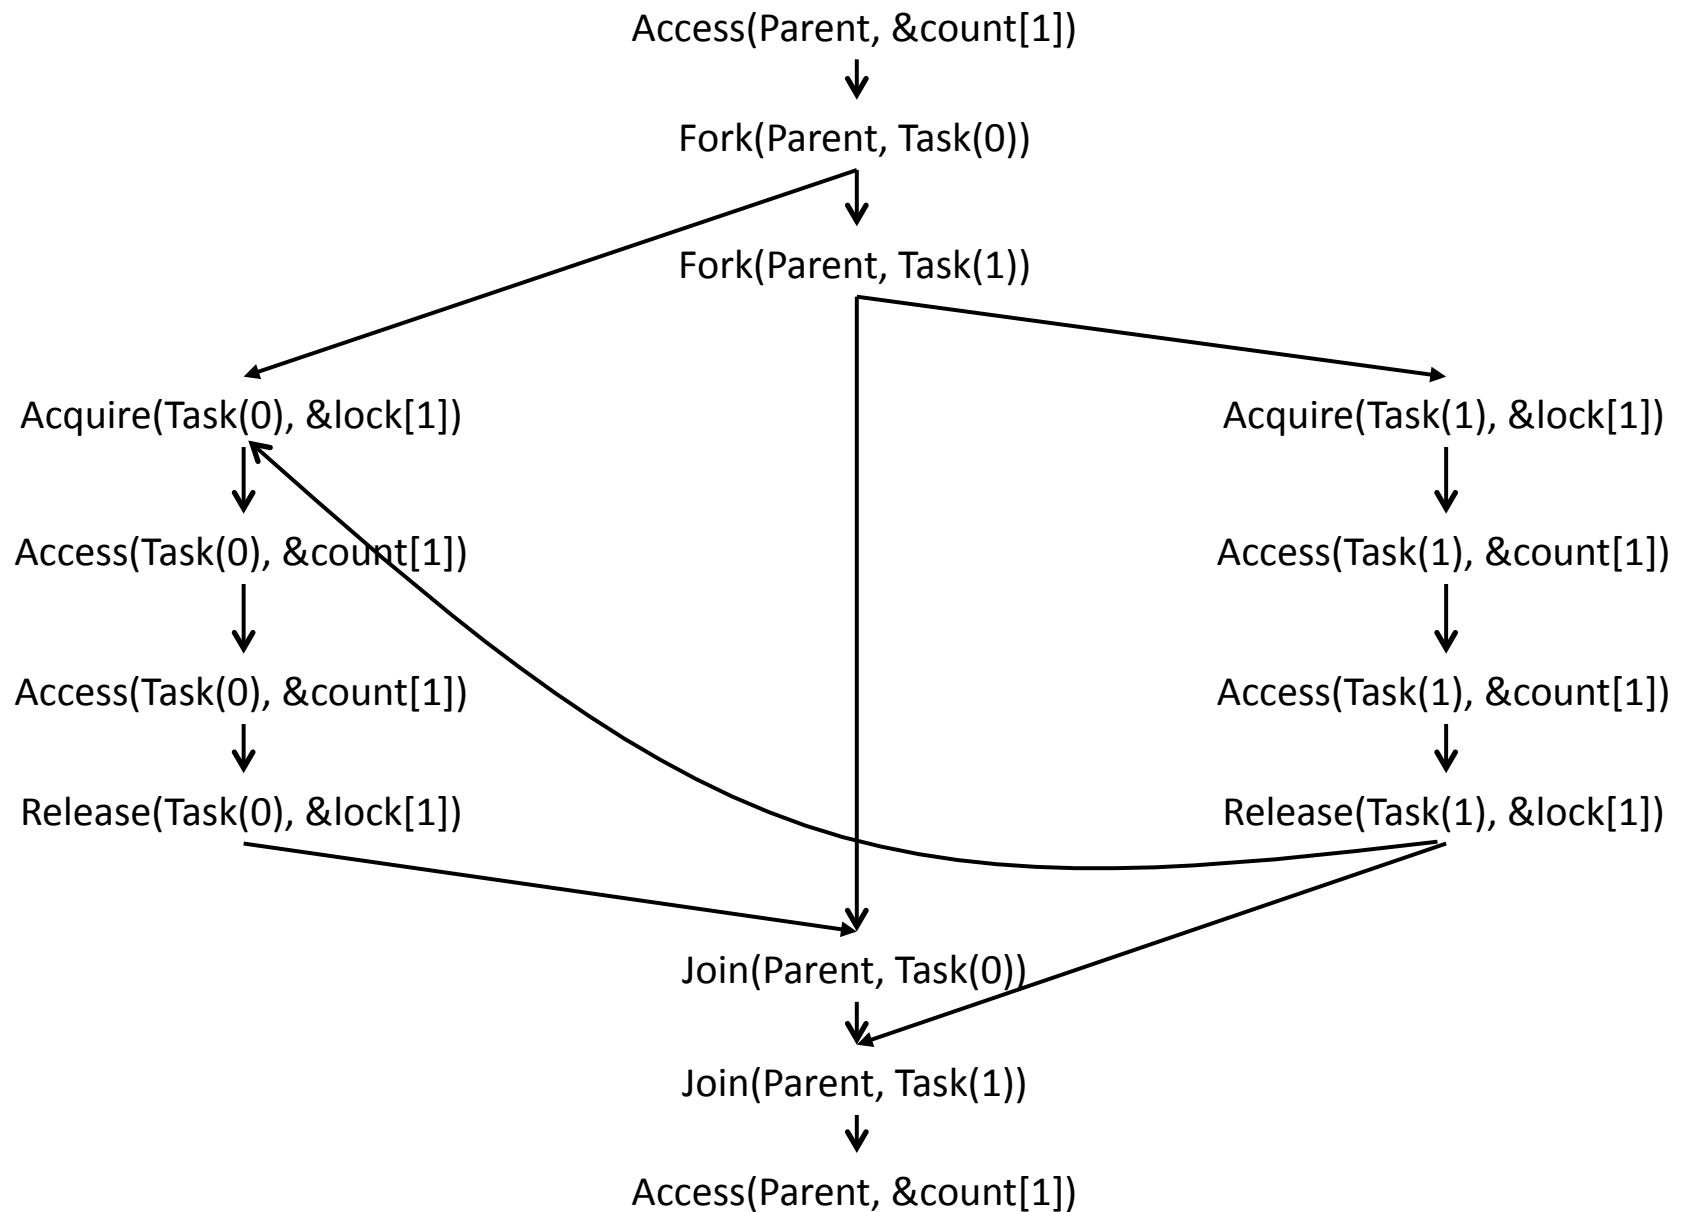
```
count[1] = t;
```

```
release(lock[1]);
```


What is a parallel execution?

- Happens-before graph: directed acyclic graph over the set of events in an execution
- Five kinds of events
 - Access(t, a): task t accessed address a
 - Fork(t, u): task t created task u
 - Join(t, u): task t waited for task u
 - Acquire(t, l): task t acquired lock l
 - Release(t, l): task t released lock l
- Three kinds of edges
 - program order: edge from an event by a particular task to subsequent event by the same task
 - fork: edge from Fork(t, u) to first event performed by task u
 - join: edge from last event performed by task u to Join(t, u)
 - release-acquire: edge from Release(t, l) to subsequent Acquire(u, l)





Vector-clock algorithm extended

- Vector clock: an array of integers indexed by the set of tasks
- For each task t , maintain a vector clock $C(t)$
 - each clock in $C(t)$ initialized to 0
- For each address a , maintain a vector clock $X(a)$
 - each clock in $X(a)$ initialized to 0
- For each lock l , maintain a vector clock $S(l)$
 - each clock in $S(l)$ initialized to 0

Vector-clock operations extended

- Task t executes an event
 - increment $C(t)[t]$ by one
- Task t forks task u
 - initialize $C(u)$ to $C(t)$
- Task t joins with task u
 - update $C(t)$ to $\max(C(t), C(u))$
- Task t accesses address a
 - data race unless $X(a) < C(t)$
 - update $X(a)$ to $C(t)$
- Task t acquires lock l
 - update $C(t)$ to $\max(C(t), S(l))$
- Task t releases lock l
 - update $S(l)$ to $C(t)$

	C(Parent)	C(Task(0))	C(Task(1))	S(&lock[1])	X(&count[1])
	[0, 0, 0]	[0, 0, 0]	[0, 0, 0]	[0, 0, 0]	[0, 0, 0]
Access(Parent, &count[1])	[1, 0, 0]				[1, 0, 0]
Fork(Parent, Task(0))	[2, 0, 0]	[2, 0, 0]			
Fork(Parent, Task(1))	[3, 0, 0]		[3, 0, 0]		
Acquire(Task(0), &lock[1])		[2, 1, 0]			
Access(Task(0), &count[1])		[2, 2, 0]			[2, 2, 0]
Access(Task(0), &count[1])		[2, 3, 0]			[2, 3, 0]
Release(Task(0), &lock[1])		[2, 4, 0]		[2, 4, 0]	
Acquire(Task(1), &lock[1])			[3, 4, 1]		
Access(Task(1), &count[1])			[3, 4, 2]		[3, 4, 2]
Access(Task(1), &count[1])			[3, 4, 3]		[3, 4, 3]
Release(Task(1), &lock[1])			[3, 4, 4]	[3, 4, 4]	
Join(Parent, Task(0))	[4, 4, 0]				
Join(Parent, Task(1))	[5, 4, 4]				
Access(Parent, &count[1])	[6, 4, 4]				[6, 4, 4]